

Efficient Formal Methods for the Synthesis of Concurrent Programs²

Paul C. Attie³

MIT Laboratory for Computer Science
Cambridge, MA
(attie@theory.lcs.mit.edu)

Introduction and background

The objective of this research is to produce useful, low-cost methods for developing correct concurrent programs from formal specifications. In particular, we address the design and verification of the synchronization and communication portions of such programs. Often, this portion can be implemented using a fixed, finite amount of synchronization related data, i.e., it is “finite-state.” Nevertheless, even when each program component contains only one bit of synchronization related data, the number of possible global synchronization states for K components is about 2^K , in general. Because of this “state-explosion” phenomenon, the manual verification of large concurrent programs typically requires lengthy, and therefore error-prone, proofs. Using a theorem prover increases reliability, but requires extensive formal labor to axiomatize and solve verification problems. Automatic verification methods (such as reachability analysis and temporal logic model checking) use state-space exploration to decide if a program satisfies its specification, and are therefore also subject to state-explosion. To date, proposed techniques for ameliorating state-explosion either require significant manual labor, or work well only when the program is highly symmetric and regular (e.g., many functionally similar components connected in similar ways).

To overcome these drawbacks, we advocate the *synthesis of programs from specifications*. This approach performs the refinement from specifications to programs automatically. Thus, the amount of formal labor is reduced to writing a formal specification and applying the appropriate synthesis step at each stage of the derivation. While nontrivial, writing a formal specification is necessary in any methodology that guarantees correctness.

Our approach

Previous synthesis methods relied on some form of exhaustive search of the program’s state-space. Hence, state-explosion has been a major obstacle to the application of synthesis to problems of realistic size. Our approach avoids exhaustive state-space search (and its exponential complexity) by analyzing interactions among every pair of component processes in the program separately, rather than looking at all processes at once. For every pair of directly interacting processes, we represent their interaction explicitly and separately from all the other interactions in the program. We start with a specification which describes all the pairwise interactions and we

automatically synthesize a “pair-program” for each such interaction. We then “compose” the pair-programs together to produce the final large program. The final step refines this program (which uses shared memory) to a message passing model. To date, we have developed:

- A method for synthesizing fault-tolerant pair-programs [AAE98].
- An efficient method for composing pair-programs into a large (shared memory) program [AE98, Att99a]. This method allows all of the pair-programs to be different.
- A methodology for refining liveness properties [Att99b]. This forms the basis for refining the large shared memory program into a distributed program, which is a next step of our research.

Our approach has synthesized programs for resource contention problems such as K -process mutual exclusion and K -process dining philosophers, for arbitrary $K \geq 2$ [AE98], and also two-phase commit [Att99a].

Concluding remarks and future work

We aim to develop a methodology that can synthesize realistic concurrent programs. Our future work will increase the scope of applicability of our method, and will address non-functional properties such as fault-tolerance, performance, and distribution.

References

- [AAE98] A. Arora, P. C. Attie, and E. A. Emerson. Synthesis of fault-tolerant concurrent programs. In *7th Annual ACM Symposium on the Principles of Distributed Computing*, pages 173 – 182, June 1998.
- [AE98] P. C. Attie and E. A. Emerson. Synthesis of concurrent systems with many similar processes. *ACM TOPLAS*, 20(1):51–115, January 1998.
- [Att99a] P. C. Attie. Synthesis of large concurrent programs via pairwise composition. In *CONCUR’99: 10th International Conference on Concurrency Theory*, number 1664 in LNCS. Springer-Verlag, August 1999.
- [Att99b] P. C. Attie. Liveness-preserving simulation relations. In *18th Annual ACM Symposium on the Principles of Distributed Computing*, pages 63 – 72, May 1999.

Relational Programs⁴

Farokh B. Bastani

University of Texas at Dallas
Richardson, TX 77083
(bastani@utdallas.edu)

Motivation

Computers are being used to automate critical services, including manufacturing systems, transportation, etc. The software for these systems is becoming very complex due to their growing sophistication. For these critical applications, it is

²This work has been supported by NSF CAREER Grant CCR-9702616 and AFOSR Grant F49620-96-1-0221.

³On leave from the School of Computer Science, Florida International University.

⁴Part of this work has been supported by NSF grant CCR-9803993 and CCR-9900922.

necessary to be able not only to *achieve* high quality but also to rigorously *demonstrate* that high quality has in fact been achieved.

Advances in software development methods, such as continuous process improvement methods, sophisticated tools, and rigorous techniques, can reduce the number of faults injected into the system during the development process. However, these methods cannot prevent or detect specification faults. One approach that is used to facilitate prevention as well as detection of these faults is to decompose the requirements specification into more manageable portions. Various decomposition methods have been proposed to simplify requirements analysis and enhance software quality. However, these methods do not necessarily enable the demonstration of high quality. For complex software systems, one way of achieving "assessability" is the divide-and-conquer approach, i.e., to be able to infer the properties of the system from those of its components. However, this inference is not always possible for arbitrary decompositions due to the difficulty of determining the operational profile for deeply nested components. Hence, after the implementation phase, we are still left with the task of determining the reliability of one complex monolithic program.

IDEAL Components

We have identified a class of software decompositions that allows various system properties (such as reliability, stability, and safety) to be inferred from the corresponding properties of the components in the system. Each component in this class is *independently developable*, i.e., it can be designed and implemented independently of the other components in the system. In addition, each component is *end-user assessable*, i.e., it can be tested or verified by the end-user independently of any other component. We refer to these components as IDEAL (Independently Developable End-user Assessable Logical) software components.

The composition of IDEAL components generally requires the components to be *relational*, i.e., the components must return all possible outputs for an input rather than just one output. The advantages of IDEAL components are as follows:

1. **Facilitates high quality software.** The global correctness properties of the system are preserved by formal composition and no further proofs are needed.
2. **Formal system integration.** Well-defined mathematical operations, i.e., union and intersection, are used to automatically compose the components together to form the system.
3. **High Safety Assurance.** The composition procedure provides strong guarantees that the safety-critical objectives will never be compromised [1].
4. **Enables fault isolation and confinement.** Faulty components can be identified and isolated quickly by analyzing the system output. Repair/recovery actions are also confined within each component separately.
5. **Supports multi-paradigm implementation.** Each

IDEAL component can be implemented using the technology that is most suitable for it without compromising system safety and reliability.

Summary

The approach has been applied to several examples drawn from the literature and one significant application based on a specification provided by Sandia National Labs. [3]. These case-studies have shown that relational programs are useful for process-control applications, communication protocols, coordination protocols, and pruning of recursive search spaces that correspond to recursive data structures. Relational programs seem to be less useful for reactive discrete systems, including telecommunication programs. However, in all these cases, the larger class of IDEAL components has proved to be very practical and useful.

References

- [1] F.B. Bastani, "Relational programs: Architecture for robust process-control programs," *Annals SE*, 1999.
- [2] F.B. Bastani, V.L. Winter, and I.-L. Yen, "Dependability of relational safety-critical programs," *Proc. ISSRE'99 FastAbstracts*, Nov. 1999.
- [3] F.B. Bastani, V. Reddy, P. Srigiriraju, and I.-L. Yen, "A relational program architecture for the Bay Area Rapid Transit system," *Proc. HIS'99*, Nov. 1999.

Reducing the Cost of Regression Testing

David Binkley
Loyola College in Maryland
(binkley@cs.loyola.edu)

The main objective of this research is to improve and invent techniques for automatic regression testing. Regression testing, which is an expensive but necessary part of software maintenance and evolution, is used to provide confidence that a modification to a program is correct. One negative consequence of regression testing's expense is that it is sometimes not performed because a maintainer "knows" that a change is correct. This often has disastrous results. The initial focus of this research is a series of experiments with a new technique for reducing the cost of (and thereby hopefully increasing the use of) regression testing.

The technique, which uses a combination of program slices to extract semantically meaningful parts of a program, produces a program called **differences** that captures the semantic change between a previously tested program and a changed version of this program; it should be more efficient to test **differences**, because it omits unchanged computations. This research should assist in the technology transfer of a rather theoretical technique to industry.

An implementation of **differences** was completed this past summer. The implementation is built on top of Code Surfer (a code analysis and slicing tool for C programs). The first