

Wait-free Byzantine Agreement

Paul Attie

College of Computer Science
Northeastern University

Abstract

We study the problem of devising wait-free consensus protocols that tolerate Byzantine failures. We show that wait-free Byzantine agreement cannot be achieved even in the presence of a single faulty (i.e., Byzantine) process. We further show that wait-free weak Byzantine agreement cannot be achieved if shared objects are “resettable,” i.e., can be reset to an initial state from any of their reachable states by an appropriate sequence of operations. Resettable objects are preferable to non-resettable objects, since they can be reused more easily. We introduce a non-resettable object, a “sticky register,” which is a straightforward generalization of the `sticky` type of [Jay93], and use it to devise a solution to wait-free weak Byzantine agreement that tolerates any number of faulty processes. Our impossibility results hold even if we can limit any process (including Byzantine processes) to invoking objects in a particular fixed set (i.e., impose a particular “access pattern”). Finally, we study the problem of *repetitive* wait-free weak Byzantine agreement: how to solve a sequence of successive instances of agreement. A straightforward solution uses an unbounded number of sticky registers—one for each instance of agreement. We show how to bound the number of registers needed by using failure detectors and weakening the wait-freedom property of the solution.

1 Introduction

Wait-freedom has been studied [Her91] extensively in recent years. A principal advantage of wait-free protocols is that the progress of each participating process is independent of the execution speeds of the other processes. In particular, a process can make progress even if the speed of all other processes is zero, i.e., if all other processes crash. Thus, a wait-free protocol with n participating processes can tolerate up to $n - 1$ crash failures. Little work has been done, however, on wait-free protocols that tolerate more severe types of faults, e.g., Byzantine faults. In this paper, we study the problem of devising wait-free shared memory consensus protocols that tolerate Byzantine faults. We show that wait-free Byzantine agreement cannot be achieved in the presence of only one Byzantine process. Our proof is straightforward, and exploits the validity condition for Byzantine agreement. We also show that wait-free weak Byzantine agreement can be achieved, but only by using “non-resettable” shared objects, that is, objects whose state cannot be reset to an initial state (by any sequence of permissible operations). With resettable objects, even a single Byzantine process is enough to prevent any solution to wait-free weak Byzantine agreement. Our impossibility results hold even if we can limit any process (including Byzantine processes) to invoking objects in a particular fixed set (i.e., impose a particular “access pattern”). Although this “limited access” model restricts the amount of damage a Byzantine process can do, the restriction is not powerful enough to enable a solution using resettable objects.

Using a non-resettable object (a sticky register — a generalization of Jayanti’s `sticky` type [Jay93] to more than two values) we give a protocol for wait-free weak Byzantine agreement. Non-resettable objects can serve for only one instance of agreement. To accommodate multiple instances (i.e., to solve the problem of *repetitive* Byzantine agreement) we need one sticky register for each instance of agreement. Because the amount of memory thus used grows unboundedly, this is not practical unless we can recover sticky registers that are “no longer needed.” To effect such recovery, we make the reasonable assumption that nonfaulty processes progress at some non-zero speed, and use a failure detector to exclude faulty processes from consideration. We can now recover all registers used for agreement instances that the slowest nonfaulty process has already participated in.

The rest of the paper is organized as follows. Section 2 defines our model of computation. Section 3 defines the wait-free Byzantine agreement problem and its weak variant. Sections 4 and 5 present our

impossibility results for wait-free Byzantine agreement and wait-free weak Byzantine agreement, respectively. Section 6 gives a solution to wait-free weak Byzantine agreement using a single sticky register. Section 7 studies the problem of repetitive wait-free weak Byzantine agreement, and Section 8 presents our conclusions and ideas for further work.

2 Model of Computation

A *concurrent program* $P = (P_1 \parallel \dots \parallel P_n, O_1, \dots, O_m, \mathcal{A})$ consists of n sequential, nondeterministic *processes* that execute concurrently and interact by performing operations on the *shared data objects* O_1, \dots, O_m . $\mathcal{A} \subseteq \{P_1, \dots, P_n\} \times \{O_1, \dots, O_m\}$ is an *access pattern*: P_i can perform operations on O_ℓ only if $(P_i, O_\ell) \in \mathcal{A}$. This access restriction remains in force even if P_i is Byzantine. Note that the access restriction is purely spatial: it only states which objects a process can invoke operations on. It says nothing about the temporal aspects of a process's behavior. An *operation* is a pair consisting of an invocation and the next matching response. A *step* of a process P_i is either an internal transition of P_i , or an operation invocation by P_i , or an operation response to P_i . Processes and objects are asynchronous state-machines: there are no upper or lower bounds on the time it takes a process to execute a step or an object to respond to an invocation. Each process has a *local state* that is not visible to the other processes, and each shared object has a local state that is not visible to the processes except via the response of that object to invocations. We designate a subset of the local state space of each process (object) as the set of *initial states* of that process (object) respectively. We assume that all shared objects are reliable (i.e., they eventually respond to any invocation) and linearizable [HW90]. [MRL98] introduces the correctness criterion of *Byznearizability*, which extends linearizability to a Byzantine environment. Our impossibility proofs also apply if this stronger correctness property of shared objects is assumed.

A *global state* is a tuple consisting of the local states of every process and shared object, and an *initial global state* is a tuple consisting of initial local states for every process and shared object. An *execution* is an alternating sequence of global states and process steps that starts in an initial global state. Hence, we model concurrency by nondeterministic interleaving. Also, we do not assume any form of fair scheduling. A *maximal execution* is an execution that is either infinite, or ends in a state in which no nonfaulty process can take a step. A *reachable state* is a state lying on some execution.

2.1 Notation and Technical Definitions

If α is an execution, then $\alpha \uparrow P_i$ is the result of removing from α all steps of processes other than P_i . Also, $\alpha \uparrow O$ is the result of removing from α all steps that are not invocations of operations on O or responses by O . This extends to a set of objects in the obvious way: $\alpha \uparrow (O_1, \dots, O_\ell)$ results from removing all invocations of or responses by some object not in O_1, \dots, O_ℓ . Finally, these notations can be nested, so that $(\alpha \uparrow P) \uparrow O$ denotes the result of removing from α all operations that are not invocations of O by P or responses to such invocations. We use $O.\text{op}(\text{arglist})$ to denote an operation **op** (with arguments *arglist*) on object O , **invoke**($O.\text{op}(\text{arglist})$) to denote an invocation of this operation, and **response**($O.\text{op}(\text{arglist})$) to denote the matching response. If s is a global state, then $s(P_i)$ is P_i 's local state in s , and $s(O)$ is object O 's local state in s . If s and t are two global states, we say $s \stackrel{i}{\sim} t$ iff $s(P_i) = t(P_i)$ and $s(O) = t(O)$ for all objects O that P_i has access to. For brevity, the constructions in this paper are expressed using pseudo-code. It is straightforward to translate this into a formal notation, e.g., I/O automata [LT88].

3 The Wait-free Byzantine Agreement problem

We propose the wait-free Byzantine agreement problem as follows. Each process starts with an initial value from a fixed set V , and each nonfaulty processes eventually decides on some value in V . The correctness conditions are as follows (the agreement and validity conditions are taken from [Lyn96, chapter 6]):

Agreement No two nonfaulty processes decide on different values.

Validity If all non-faulty processes start with the same initial value $v \in V$, then v is the only possible decision value for a nonfaulty process.

Termination If P_i is nonfaulty, then P_i decides in every maximal execution.

No-conditional-waiting If P_i is a nonfaulty process and s is a reachable state in which P_i has not decided, then there exists an execution α starting in s such that:

1. α consists only of steps of P_i .
2. P_i decides in α .

Uniform-initial-state Every combination of initial values for each process and initial local state for each shared object is a possible initial global state (of any solution).

The **Termination** condition rules out busy waiting, and the **No-conditional-waiting** condition rules out conditional waiting. Together, these impose the wait-freedom requirement (cf. the definition in [Her91, p. 129]). We also define the wait-free weak Byzantine agreement problem, in which the Validity condition is replaced by the Weak Validity condition:

Weak Validity If there are no faulty processes and all processes start with the same initial value $v \in V$, then v is the only possible decision value.

We impose the uniform initial state condition in order to rule out trivial solutions in which each process has a “shared” object that only it can access, and whose initial value gives the correct decision value. In the sequel, we assume without loss of generality that $\{0, 1\} \subseteq V$.

4 Impossibility of Wait-free Byzantine Agreement

In this section, we show that there is no solution to wait-free Byzantine agreement, even if a single process exhibits Byzantine behavior. Our first proposition shows that, when a process P_i runs alone and decides before other processes take any steps, then it must decide its own initial value. Intuitively, this follows because P_i ’s decision must be independent of the initial values of the other processes, and so, to satisfy validity, P_i must decide on its own initial value.

Proposition 1 *Let $P = (P_1 \parallel \dots \parallel P_n, O_1, \dots, O_m, \mathcal{A})$ be a protocol for wait-free Byzantine agreement. In any execution of P in which P_i , $i \in \{1..n\}$, runs alone and decides before any other process takes a step, P_i must decide its own initial value.*

Proof. Let s be an arbitrary global initial state, and let α_s be an arbitrary execution starting in s in which P_i first runs alone until it decides. Let the initial value of P_i in s_i be v_i . Let t be an arbitrary global initial state which is identical to s except that all processes have initial value v_i (thus all shared objects have the same initial local states in s as in t). Starting in state t , it is possible for P_i to execute the same steps that it executes in α_s , since P_i ’s local state and all object states are the same in s as in t . Let α_t be the resulting execution. By validity, P_i decides v_i in α_t . Hence, P_i also decides v_i in α_s since it executes the same steps in both executions. Since α_s was chosen arbitrarily, the proposition follows. \square

We first establish the impossibility result for three processes. The generalization to n processes is easily done using the no-conditional-waiting and uniform-initial-state requirements.

Theorem 2 *There is no solution to the wait-free Byzantine agreement problem for three processes if one of them can be faulty.*

Proof. By contradiction. Suppose a solution $(P_1 \parallel P_2 \parallel P_3, O_1, \dots, O_m, \mathcal{A})$ exists. Let α_1 be the following execution of this solution: P_1, P_2 are nonfaulty and have initial values 1, 0 respectively. First, P_1 runs alone until it decides. By proposition 1, P_1 decides 1. Then, P_2 runs alone until it decides. By the agreement condition, P_2 also decides 1.

Now consider execution α_2 which is as follows. P_1 is faulty, P_2 is nonfaulty and has the same initial state as in α_1 (and therefore has initial value 0), and P_3 is nonfaulty and has initial value 0. Also, the shared objects have the same initial states as in α_1 . First, P_1 runs alone, and executes exactly the same steps as it does in α_1 . This is possible since the shared objects have the same initial values and P_1 is faulty (and can therefore deviate from the protocol). Next, P_2 runs alone until it decides. Since the shared objects have the same values in α_1 and α_2 when P_2 starts executing, and P_2 has the same initial state, it is possible for P_2 to execute exactly the same sequence of steps that it executed in α_1 . Let this then be the sequence of steps that P_2 executes in α_2 . Hence P_2 decides 1 in α_2 , since it decides 1 in α_1 . But this contradicts the validity condition, and we are done. \square

Theorem 3 *There is no solution to the wait-free Byzantine agreement problem for $n \geq 2$ processes if one of them can be faulty.*

Proof. For $n = 2$, the proof is essentially the same as that for Theorem 6.27 in [Lyn96], and is omitted. For $n \geq 3$, we assume a solution exists and derive a contradiction exactly as in the proof of Theorem 2. From the wait-freedom condition and the uniform initial state condition, it is clear that the executions α_1, α_2 (with α_1 's initial state extended arbitrarily to P_4, \dots, P_n , and α_2 's initial state chosen so that P_4, \dots, P_n also have initial value 0) constructed in the aforementioned proof are executions of a solution for any $n \geq 3$. \square

5 Impossibility of Wait-free Weak Byzantine Agreement using Resettable Objects

While wait-free Byzantine agreement is not achievable, weak wait-free Byzantine agreement can be solved, but only by using shared objects which cannot be reset to an initial state. We define:

Definition 1 *An object is resettable if, for each of its reachable states, there exists a sequence of operations that takes the object back to some initial state. We call such a sequence of operations a reset sequence.*

Note that our definition allows an object to be reset to several initial states from a given reachable state s . The object cannot however, be reset to an initial state that is not reachable from s . An extreme case is when every initial state is reachable from s , and so the object could be reset to any initial state (from s).

Resettable objects are more desirable than non-resettable objects, because they can be easily reused (e.g., for successive instances of agreement). A non-resettable object can only be “reused” by reclaiming the memory that it occupies, i.e., destroying it, and dynamically creating a new instance of the same object type. We examine these issues further in Section 7.

In this section, we show that there is no solution to wait-free weak Byzantine agreement using only resettable objects, and in the presence of a single Byzantine process. In the next section, we exhibit a solution that uses a non-resettable object: a sticky register.

Proposition 4 *Let $P = (P_1 \parallel \dots \parallel P_n, O_1, \dots, O_m, \mathcal{A})$ be a protocol for wait-free weak Byzantine agreement. In any execution of P in which $P_i, i \in \{1..n\}$, runs alone and decides before any other process takes a single step, P_i must decide its own initial value.*

Proof. In the proof of Proposition 1, no processes are assumed faulty. Hence, the proof still applies if we assume that all processes are nonfaulty, and we replace the use of validity by the use of weak validity. \square

As before, we present the impossibility result for three processes, and easily generalize it to n processes using the no-conditional-waiting and uniform-initial-state requirements.

Theorem 5 *There is no solution to the wait-free weak Byzantine agreement problem for three processes if one of them can be faulty and all shared objects are resettable.*

Proof. By contradiction. Suppose a solution $P = (P_1 \parallel P_2 \parallel P_3, O_1, \dots, O_m, \mathcal{A})$ exists. Let $\mathcal{O}_{ij} \subseteq \{O_1, \dots, O_m\}$ be the set of objects that can be accessed by P_i and P_j , $i, j \in \{1, 2, 3\}, i \neq j$, and $\mathcal{O}_{123} \subseteq \{O_1, \dots, O_m\}$ be the set of objects that can be accessed by P_1, P_2 , and P_3 , all according to the access pattern \mathcal{A} . We construct two executions α_1 and α_2 of P as follows (Figure 1 shows how α_1 and α_2 are “alternately” constructed from each other in a well-defined manner).

The first part of α_1 is as follows. P_1 is faulty, P_3 has initial value 1, and \mathcal{O}_{23} is in some initial state \mathcal{V}_{23} ¹. First, P_2 runs alone until it decides. Without loss of generality, suppose P_2 decides 1. Next, P_1 invokes reset sequences for all objects in \mathcal{O}_{13} and \mathcal{O}_{123} . Let the resulting initial states of $\mathcal{O}_{13}, \mathcal{O}_{123}$ be $\mathcal{V}_{13}, \mathcal{V}_{123}$ respectively.

α_2 is as follows. P_2 is faulty, P_1 is nonfaulty and has initial value 0, and P_3 is nonfaulty and has the same initial state as in α_1 (and therefore has initial value 1). The initial states of $\mathcal{O}_{13}, \mathcal{O}_{23}, \mathcal{O}_{123}$ are $\mathcal{V}_{13}, \mathcal{V}_{23}, \mathcal{V}_{123}$ respectively (by the uniform-initial-state requirement, these initial conditions must exist for some execution). First, P_1 runs alone until it decides. By Proposition 4, P_1 decides 0. Next, P_2 executes $(\alpha_1 \uparrow P_2) \uparrow \mathcal{O}_{23}$ (since P_2 is faulty, it can execute any sequence of operations on objects it has access to). Let the global state at this point be $\alpha_2.t$. Finally, P_3 runs alone until it decides. By agreement, P_3 must decide 0, since P_1 decided 0, and P_1, P_3 are nonfaulty.

The last part of α_1 is as follows. P_1 executes $(\alpha_2 \uparrow P_1) \uparrow (\mathcal{O}_{13} \cup \mathcal{O}_{123})$ —since P_1 is faulty, it can execute any sequence of operations on objects it has access to. Let the global state at this point be $\alpha_1.s$. Finally, P_3 runs alone until it decides. By agreement, P_3 must decide 1, since P_2 decided 1, and P_2, P_3 are nonfaulty.

We now establish $\alpha_1.s \stackrel{3}{\sim} \alpha_2.t$. In $\alpha_1.s$, the state of \mathcal{O}_{23} results from the sequence $(\alpha_1 \uparrow P_2) \uparrow \mathcal{O}_{23}$ applied to the initial state \mathcal{V}_{23} of \mathcal{O}_{23} (by definition of \uparrow). In $\alpha_2.t$, the state of \mathcal{O}_{23} results from the same sequence, namely $(\alpha_1 \uparrow P_2) \uparrow \mathcal{O}_{23}$, applied to the same initial state \mathcal{V}_{23} (by construction of α_2). Hence $\alpha_1.s(\mathcal{O}_{23}) = \alpha_2.t(\mathcal{O}_{23})$.

In $\alpha_1.s$, the states of $\mathcal{O}_{13}, \mathcal{O}_{123}$ result from the sequence $(\alpha_2 \uparrow P_1) \uparrow (\mathcal{O}_{13} \cup \mathcal{O}_{123})$ applied to the states $\mathcal{V}_{13}, \mathcal{V}_{123}$ respectively (because of the reset sequences invoked by P_1 in α_1 , the actual initial states of $\mathcal{O}_{13}, \mathcal{O}_{123}$ don’t matter). In $\alpha_2.t$, the states of $\mathcal{O}_{13}, \mathcal{O}_{123}$ also result from the sequence $(\alpha_2 \uparrow P_1) \uparrow (\mathcal{O}_{13} \cup \mathcal{O}_{123})$ applied to the (initial) states $\mathcal{V}_{13}, \mathcal{V}_{123}$ respectively (by definition of \uparrow). Hence we conclude $\alpha_1.s(\mathcal{O}_{13}) = \alpha_2.t(\mathcal{O}_{13})$, $\alpha_1.s(\mathcal{O}_{123}) = \alpha_2.t(\mathcal{O}_{123})$. Finally, the local state of P_3 in $\alpha_1.s$ is its initial state in α_1 , since P_3 takes no steps in the prefix of α_1 up to $\alpha_1.s$. Likewise the local state of P_3 in $\alpha_2.t$ is its initial state in α_2 . But these initial states are the same, by construction of α_2 . Hence $\alpha_1.s(P_3) = \alpha_2.t(P_3)$. Since P_3 ’s state and the states of all objects that P_3 can access are the same in $\alpha_1.s$ as in $\alpha_2.t$, we conclude $\alpha_1.s \stackrel{3}{\sim} \alpha_2.t$.

In α_1 , P_3 runs alone from $\alpha_1.s$ until it decides. In α_2 , P_3 runs alone from $\alpha_2.t$ until it decides. Since $\alpha_1.s \stackrel{3}{\sim} \alpha_2.t$, and P_3 runs alone from $\alpha_1.s, \alpha_2.t$ until it decides, it is possible for P_3 to execute the same sequence (i.e., $\alpha_1 \uparrow P_3$) of steps in α_2 as in α_1 . Let this then be the sequence that P_3 actually executes in α_2 . Hence, P_3 must decide the same value in α_2 as in α_1 . But we showed above that P_3 decides 1 in α_1 and 0 in α_2 . Hence the desired contradiction. \square

Theorem 6 *There is no solution to the wait-free weak Byzantine agreement problem for $n \geq 3$ processes if one of them can be faulty and all shared objects are resettable.*

Proof. We assume a solution exists and derive a contradiction exactly as in the proof of Theorem 5. From the wait-freedom condition and the uniform initial state condition, it is clear that the executions α_1, α_2 (with their initial states extended arbitrarily to P_4, \dots, P_n) constructed in the aforementioned proof are executions of a solution for any $n \geq 3$. \square

We note that if an access pattern is not imposed, then the proof of Theorem 5 becomes trivial: let P_1 be nonfaulty and run alone until it decides; then let P_2 be faulty and reset all objects in the system; finally

¹For brevity, we discuss the sets of objects $\mathcal{O}_{ij}, \mathcal{O}_{123}$ as if they were single objects. For example, \mathcal{V}_{23} is actually a set of initial object states.

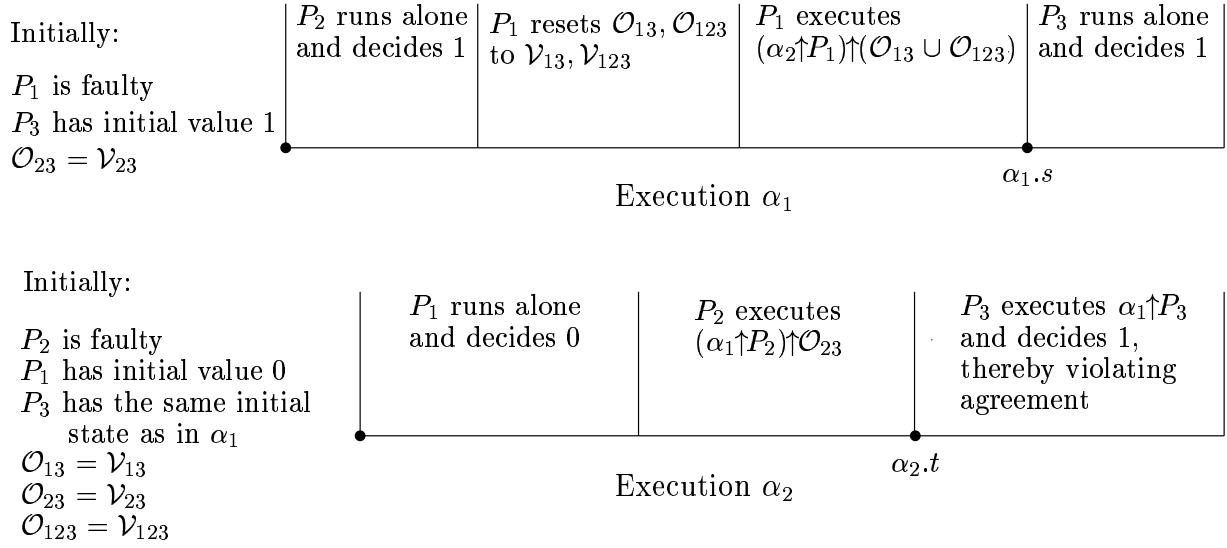


Figure 1: Executions α_1 and α_2 in the proof of Theorem 5.

let P_3 be nonfaulty and run alone until it decides. Invoking Proposition 4 then allows us to conclude that P_1 and P_3 must both decide on their initial values, since P_3 's view of the execution is indistinguishable from its view of some execution in which it runs alone and decide before any other process takes a step.

6 Wait-free Weak Byzantine Agreement Using One Sticky Register

A sticky register is a register whose value can be set once, and subsequently can never be changed. It is essentially the same as the **consensus** type of [Her91], and generalizes the **sticky** type of [Jay93] from a binary domain to an arbitrary domain (cf. also the sticky bit of [Plo89]). A sticky register r has initial value $\perp \notin V^2$, and supports a single operation, $r.\text{jam}(v)$ ³, that sets r 's value to v if the value is \perp and otherwise leaves r unchanged. In either case, $r.\text{jam}(v)$ returns the resulting value of r . The formal definition of $r.\text{jam}(v)$ is as follows:

```

 $r.\text{jam}(v : V)$  returns( $V$ )
  if  $r = \perp \wedge v \in V$  then  $r := v$  endif;
  return( $r$ )

```

Given a sticky register r , we have a trivial protocol $(P_1 \parallel \dots \parallel P_n, r)$ that solves wait-free weak Byzantine agreement: each P_i executes $r.\text{jam}(v_i)$, where v_i is P_i 's initial value, and decides on the value returned. The protocol for P_i , $i \in \{1..n\}$, is:

```

agree( $r$  : sticky register,  $v_i : V$ )
  invoke( $r.\text{jam}(v_i)$ );
   $d_i := \text{response}(r.\text{jam}(v_i))$ ;
  decide( $d_i$ )

```

Proposition 7 *The above protocol solves the wait-free weak Byzantine agreement problem, in the presence of any number of Byzantine processes.*

²Recall that V is the domain of possible decision values.

³We borrow Plotkin's [Plo89] terminology here.

Proof. For agreement, let v be the value returned by the first `jam` operation. By definition of `jam`, all subsequent `jam` operations also return v , even in the presence of Byzantine behavior by any number of processes. Hence all nonfaulty processes decide v . For weak validity, if all processes are nonfaulty and have the same initial value v , then the first `jam` operation must be $r.\text{jam}(v)$, by definition of the protocol. Hence, all `jam` operations return v , and so all processes decide v . Termination and no-conditional-waiting are obvious, since nonfaulty processes execute a finite amount of straight-line code that contains no tests. Finally, the sticky register r has only one possible initial state, namely \perp , and the only local state that each process has is its initial value. Hence it is clear that the protocol permits all possible initial global states, i.e., the uniform initial state requirement is satisfied. \square

It is easy to see that this protocol works because the behavior of a sticky register severely limits the amount of damage that a Byzantine process can inflict. All that a Byzantine process can do is set a sticky register to an arbitrary value in V . This is sufficient to destroy validity (and hence Theorem 3), but not to destroy agreement.

7 Repetitive Wait-free Weak Byzantine Agreement

The protocol of Section 6 uses a single sticky register to solve a single instance of wait-free weak Byzantine agreement. Thus, to solve *repetitive* agreement—where processes repeatedly engage in successive instances of agreement—we need an unbounded number of sticky registers. This is because sticky registers are not resettable, and so cannot be reused, which means that we need a new sticky register for each instance of agreement. We assume that an underlying, reliable operating system is responsible for allocating and deallocating registers. We represent the operating system as an additional process OS .

Our protocol for repetitive agreement is as follows. We number the instances of agreement with the positive integers. We also number each sticky register with the number of the agreement instance that the register is used for. Each process P_i maintains an instance variable $instance_i$. After participating in some instance of agreement, P_i increases $instance_i$ to the number of the next agreement instance that it wishes to participate in. Note that the increments of $instance_i$ need not be 1, i.e., not every nonfaulty process need participate in every agreement instance.

As mentioned above, this protocol uses an unbounded number of sticky registers. To bound the number of sticky registers needed at any time, we:

1. Reclaim registers that are “no longer needed.” A register is no longer needed when the “slowest” nonfaulty process has participated in the agreement instance that the register is used for.
2. Occasionally force a “fast” nonfaulty process to wait for more sticky registers to be allocated. This entails giving up some aspects of wait-freedom. Specifically, processes are still wait-free “relative to one another,” i.e., one process cannot delay another, but processes are not wait-free relative to the operating system (OS), since OS can delay a process by forcing it to wait for a sticky register to be allocated. We define *pseudo-wait-free* Byzantine agreement by modifying the termination and no-conditional-waiting requirements as follows:

Termination If P_i is nonfaulty, then P_i decides in every maximal execution in which OS allocates the required sticky register.

No-conditional-waiting Let P_i be a nonfaulty process and s be a reachable state in which P_i has not decided. Then, if there exists an execution starting in s in which OS allocates the required sticky register, then there also exists an execution α starting in s such that:

- (a) α consists only of steps of P_i or OS .
- (b) P_i decides in α .

By using memory freed up by reclamation to allocate new registers, we can bound the total amount of memory used for sticky registers at any time. Also, since processes probably have to wait for the operating

system at some point, the replacement of wait-freedom by pseudo-wait-freedom is unlikely to be significant from a practical point of view. The above discussion suggests a generalization of wait-freedom to “wait-freedom relative to another particular process.” We defer the formal definition and exploration of this notion to another occasion.

To reclaim registers that are no longer needed, OS periodically executes:

$$min_inst := \text{MIN}_{P_i \text{ is nonfaulty}} instance_i,$$

and recovers all sticky registers up to register number $min_inst - 1$. The main problem then, is detecting faulty processes and excluding them from consideration when computing min_inst . Unfortunately, since one of the possible behaviors of a Byzantine process is to behave “correctly,” it is not possible to detect all faulty processes. Following [MR97], we detect only those Byzantine processes whose behavior prevents progress, where in this case progress occurs when min_inst is increased, thereby allowing some more sticky registers to be reclaimed. Thus, for purposes of this section only, we redefine a faulty process as follows⁴:

$$P_i \text{ is faulty iff } \exists v, \Diamond \Box (instance_i^{OS} < v) \vee \neg \forall v, \Box (instance_i^{OS} = v \Rightarrow \Box (instance_i^{OS} \geq v)).$$

Here, $instance_i^{OS}$ is a variable local to the operating system process OS which records OS ’s most recent observation of the value of $instance_i$. Since our main concern here is that faulty processes prevent the reclamation of registers, it suffices to exclude from consideration faulty processes that *appear to be slow* to OS . Faulty processes that OS “perceives” to be fast, i.e., whose $instance_i$ variables are always increasing when sampled by OS , will not prevent register reclamation. For purposes of register reclamation, these processes can be treated by OS as if they were correct.

Our definition of faulty above considers a process faulty iff its observed instance number is either bounded by some value v , or decreases at some point. In the first case, this prevents reclamation of registers numbered higher than v . In the second case, if the observed instance numbers of two or more processes alternately show decreases in such a way that their minimum is bounded by some value v (i.e., the decreases “overlap”) this would again prevent reclamation of registers numbered higher than v . To detect faulty processes, we employ a failure detector \mathcal{D} [CT96]. Unlike the definition of failure detector given in [CT96], where each process has a failure detector module, our failure detector consists of only one module, which is accessible to OS . \mathcal{D} uses the instance numbers $instance_i^{OS}$ which OS observes. The completeness and strong accuracy properties of failure detectors given in [CT96] are now rephrased as follows:

Completeness Eventually, every process that fails is permanently suspected by OS .⁵

Strong accuracy No process is suspected by OS before it fails.

Eventual strong accuracy There is a time after which correct processes are not suspected by OS .

Figure 2 gives our solution for repetitive pseudo-wait-free weak Byzantine agreement with register reclamation. Each P_i repeatedly participates in the instance of agreement given by its $instance_i$ variable, and then sets $instance_i$ to the next value of a monotonically increasing sequence $schedule_i$, which gives the instances of agreement that P_i participates in⁶. We assume that $v_{instance_i}$ gives the initial value of P_i for the agreement instance identified by $instance_i$.

We require that \mathcal{D} be such that the following hold:

$$\neg \exists v, \Diamond \Box (min_inst < v) \tag{1}$$

$$\forall \text{ nonfaulty } P_i, \Box (min_inst \leq instance_i) \tag{2}$$

⁴Note that we use first-order linear-time temporal logic [Eme90]. \Box, \Diamond are the linear temporal logic operators “always” and “eventually,” respectively. Thus, $\Diamond \Box f$ means that there exists a state after which f holds continuously.

⁵Since there is only one failure detector module, the strong and weak completeness properties of [CT96] coincide.

⁶This “infinite” representation is for convenience only. In practice, P_i would compute “on the fly” the number of the next instance of agreement that it wishes to participate in.

```

 $P_i ::$    $k_i := 1;$ 
          $instance_i := schedule_i[k_i];$ 
         repeat forever
           wait until  $OS$  has allocated registers up to register number  $instance_i$ ;
           agree( $r_{instance_i}, v_{instance_i}$ );
            $k_i := k_i + 1;$ 
            $instance_i := schedule_i[k_i]$ 
         endrepeat

 $OS ::$  repeat forever
         forall  $i \in \{1..n\}$   $instance_i^{OS} := instance_i$  endfor;
          $min\_inst := \text{MIN}_{P_i \notin \mathcal{D}} instance_i^{OS};$ 
         reclaim all registers up to register  $min\_inst - 1$ ;
         allocate some new registers, possibly using memory freed up by reclamation
       endrepeat

```

Figure 2: Repetitive pseudo-wait-free weak Byzantine agreement with register reclamation.

Formula (1) requires that min_inst increase without bound. This enables the continuing reclamation of sticky registers. (2) requires that min_inst is never greater than the instance number of a nonfaulty process. This ensures that no register is reclaimed before it has been used by every nonfaulty process that needs it (i.e., every nonfaulty process that participates in the agreement instance which the register implements).

We show that completeness and strong accuracy of \mathcal{D} are sufficient to assure (1) and (2).

Proposition 8 *If \mathcal{D} satisfies completeness and strong accuracy, then (1) and (2) hold.*

Proof. By completeness of \mathcal{D} , every faulty process is eventually excluded from the computation of min_inst . Since $instance_i^{OS}$ increases without bound when P_i is nonfaulty, it follows that min_inst also increases without bound, and so (1) holds. By strong accuracy of \mathcal{D} , $instance_i^{OS}$ is included in the computation of min_inst for every nonfaulty P_i , and so $min_inst \leq instance_i^{OS}$. Since $instance_i$ increases monotonically, it follows that $instance_i^{OS} \leq instance_i$, is invariant (assume that $instance_i^{OS}$ is initialized to, say, 0). Hence $min_inst \leq instance_i$ is invariant, which establishes (2). \square

Strong accuracy may be difficult to achieve in practice. Suppose we use a failure detector that satisfies completeness and eventual strong accuracy. It is easy to see that (1) continues to hold, but that (2) may not, because now a nonfaulty process P_i may be incorrectly suspected, which would cause $instance_i$ to be excluded from the calculation of min_inst . However, the following does hold:

$$\forall \text{ nonfaulty } P_i, \Diamond \Box (min_inst \leq instance_i) \quad (2')$$

In other words, (2) holds from some point onwards. Hence, replacing a strongly accurate detector by an eventual strongly accurate detector results in a finite number of agreement instances not executing correctly because the needed sticky register was reclaimed too soon (the solution of Figure 2 must be modified so that invocations of **agree**($r_{instance_i}, v_{instance_i}$) return a default value if $r_{instance_i}$ has been reclaimed—this still satisfies agreement, but weak validity may be violated).

Proposition 9 *If \mathcal{D} satisfies completeness and eventual strong accuracy, then (1) and (2') hold.*

Proof. The proof of (1) is exactly the same as in the proof of Proposition 8 above. By eventual strong accuracy of \mathcal{D} , $instance_i^{OS}$ is eventually included in every computation of min_inst for every nonfaulty P_i (from some point on), and so $\Diamond \Box (min_inst \leq instance_i^{OS})$. Since $instance_i^{OS} \leq instance_i$, is invariant, $\Diamond \Box (min_inst \leq instance_i)$ holds, which establishes (2'). \square

8 Conclusions and Further Work

In this paper, we studied the problem of wait-free consensus in the presence of Byzantine failures. We showed that wait-free Byzantine agreement is not achievable, and that weak wait-free Byzantine agreement is achievable only if we use non-resettable objects. Our impossibility results hold in the presence of only a single Byzantine process, and even if we can impose an “access pattern” that, for each process, fixes the objects that the process has access to. We also gave a protocol for weak wait-free Byzantine agreement that uses non-resettable objects, and tolerates any number of Byzantine faults. Our results suggest that the usual “majority voting” considerations that result in the usual bound of at most one-third faulty processes ($n \geq 3f + 1$) do not come into play when wait-freedom is required. Finally, we studied repetitive weak wait-free Byzantine agreement, and showed that it is possible to bound the amount of memory needed by using failure detectors and weakening the wait-freedom property of the solution.

The access patterns considered in this paper either grant a process access to all of the operations defined by the type of a particular shared object, or to none of them. A more refined notion would grant different processes access to different subsets of the defined operations. Thus, some processes may only have access to a “read” operations. It would be worthwhile to investigate whether our impossibility results carry over to this more refined model. Also, the notion of wait-freedom relative to another process is worth investigating further. Finally, our solution of repetitive wait-free weak Byzantine agreement should be extended to use bounded instance numbers.

References

- [CT96] T.D. Chandra and S. Toueg. Unreliable failure detectors for reliable distributed systems. *J. ACM*, 43(4):225–267, Mar. 1996.
- [Eme90] E. A. Emerson. Temporal and modal logic. In J. Van Leeuwen, editor, *Handbook of Theoretical Computer Science*, volume B, *Formal Models and Semantics*. The MIT Press/Elsevier, Cambridge, Mass., 1990.
- [Her91] M. Herlihy. Wait-free synchronization. *ACM Trans. Program. Lang. Syst.*, 11(1):124–149, Jan. 1991.
- [HW90] M. Herlihy and J.M. Wing. Linearizability: a correctness condition for concurrent objects. *ACM Trans. Program. Lang. Syst.*, 12(3):463–492, July 1990.
- [Jay93] P. Jayanti. On the robustness of Herlihy’s hierarchy. In *12’t ACM-PODC Symposium*, Ithaca, NY, Aug. 1993.
- [LT88] N.A. Lynch and M.R. Tuttle. An introduction to input/output automata. Technical Report MIT/LCS/TM-373, MIT Laboratory for Computer Science, Boston, Mass., 1988.
- [Lyn96] N. Lynch. *Distributed Algorithms*. Morgan-Kaufmann, San Francisco, CA, 1996.
- [MR97] D. Malkhi and M. Reiter. Unreliable intrusion detection in distributed computations. In *10’t Computer Security Foundations Workshop*, Rockport, MA, June 1997.
- [MRL98] D. Malkhi, M. Reiter, and N. Lynch. A correctness condition for memory shared by Byzantine processes. Private communication, July 1998.
- [Plo89] S. Plotkin. Sticky bits and the universality of consensus. In *8’t ACM-PODC Symposium*, Edmonton, Alberta, Canada, Aug. 1989.