

# Efficiently Verifiable Sufficient Conditions for Deadlock-freedom of Large Concurrent Programs (Technical Report)

Paul C. Attie and Hana Chockler  
College of Computer Science, Northeastern University,  
Cullinane Hall, 360 Huntington Avenue,  
Boston, Massachusetts 02115.  
Email: {attie,hanac}@ccs.neu.edu

May 17, 2004

## Abstract

We present two polynomial-time algorithms for automatic verification of deadlock-freedom of large finite-state concurrent programs. We consider shared-memory concurrent programs in which a process can nondeterministically choose amongst several (enabled) actions at any step. Our algorithms are sound but incomplete: if they return a positive answer, then the program is indeed deadlock-free, while a negative answer conveys no information: the program might be deadlock free, but “fails” the test that our algorithms apply.

Our algorithms apply to programs which are expressed in a particular syntactic form, in which variables are shared between pairs of processes, and the synchronization code for each pair of interacting processes is expressed separately from synchronization code for other (even intersecting) pairs of processes. The first algorithm is an improvement of the deadlock check of [2]. Its complexity is polynomial in all its parameters, as opposed to the exponential complexity of the deadlock check of [2]. It also scales well and can be applied to very large concurrent programs. The second algorithm involves a conceptually new construction of a global wait-for graph for all processes. Its running time is also polynomial in all its parameters, and it is more discriminating than the first algorithm. The algorithms apply to systems with similar and with dissimilar processes.

We illustrate our algorithms by applying them to several examples of concurrent programs that implement mutual exclusion and priority queues. To the best of our knowledge, this is the first work that describes polynomially checkable conditions for assuring deadlock freedom of systems of similar and dissimilar processes.

## 1 Introduction

One of the important correctness properties of concurrent programs is the absence of *deadlocks*, e.g. as defined in [19]: “a set of processes is deadlocked if each process in the set is waiting for an event that only another process in the set can cause.” Most approaches to deadlock assume that the “event” that each process waits for is the release of a resource held by another process. We refer to this setting as the *resource allocation* setting. Four conditions are necessary for a deadlock to arise [5, 13]: (1)

resources can be held by at most one process, (2) processes can hold some resources while waiting to acquire several (more than 1, in general) others, (3) resources cannot be taken away from a process (no preemption), (4) a cyclical pattern of waiting amongst the involved processes. The exact pattern of waiting required to cause a deadlock depends on the specific resource model, and can be depicted in terms of a *wait-for-graph* (WFG), i.e., a graph whose edges depict the “wait-for” relationships between processes. The following models have been formulated [14]: (1) AND model: a process blocks iff one or more of the resources it has requested are unavailable, (2) OR model: a process blocks iff all of the resources it has requested are unavailable, (3) AND-OR model: a process can use any combination of AND and OR operators in specifying a resource request, (4)  $k$ -out-of- $n$ : a process requests any  $k$  resources out of a pool of  $n$  resources. For the AND-model, deadlock arises if the WFG contains a cycle. For the OR-model, deadlock arises if the WFG contains a knot, i.e., a set of processes each of which can reach exactly all the others by traversing wait-for edges. To our knowledge, no graph-theoretic construct characterizing deadlock in the AND-OR or the  $k$ -out-of- $n$  models is known [14].

In this paper, we address a version of the deadlock problem that is more general than the resource-based model. We consider the deadlock problem in the case that the event which each process waits for is the truthification of a predicate over shared state. Thus, we deal with a shared variables model of concurrency. However, our approach is applicable in principle to other models such as message passing or shared events. We exploit the representation of concurrent programs in a form where the synchronization between processes can be factored out, so that the synchronization code for each pair of interacting processes is expressed separately from that for other pairs, even for two pairs that have a process in common. This “pairwise” representation was introduced in [2], where it was used to synthesize programs efficiently from CTL specifications.

Traditionally, three approaches to dealing with deadlock have been investigated: (1) deadlock detection and recovery: since a deadlock is stable, by definition, it can be detected and then broken, e.g., by the preemption, rollback, or termination of an involved process. (2) deadlock avoidance: avert the occurrence of a deadlock by taking appropriate action. Deadlock avoidance algorithms have been devised for the resource-based formulation of deadlock [19], (3) deadlock prevention: prevent a deadlock from arising by design. In particular, attempt to negate one of the four conditions mentioned above for the occurrence of deadlock. As Tanenbaum [19] observes, attempting to negate any of the first three conditions is usually impractical, and so we are left with condition (4): a cyclical pattern of waiting.

**Related work** As shown in [15], deadlock-freedom analysis is NP-complete even for simple systems. Therefore, the research in this area concentrated either on the search for efficiently checkable sufficient conditions for deadlock-freedom, or on improving the complexity of the check in some special cases.

There are two approaches in the literature for preventing the cyclical waiting in concurrent systems. One approach is presented in [2] for systems of similar processes. Essentially, this approach concentrates on checking sufficient conditions for assuring deadlock freedom. The conditions should be efficiently checkable. Attie and Emerson [2] formulate a condition, checkable in subsystems consisting of  $n + 2$  processes, where  $n$  is the maximal branching degree of a state-node in a state-transition graph that represents the behavior of processes (essentially,  $n$  reflects the degree of non-determinism of a single process in a state). This condition implies absence of deadlocks in the whole system. The size of the subsystems that have to be checked is exponential in  $n$ , which makes checking the condition infeasible for systems that allow a high degree of non-determinism. Most model checking algorithms can be applied to verifying deadlock freedom. The main impediment is state-explosion. Some approaches to ameliorating

state-explosion are by using a partial order instead of an interleaving model [10, 12, 11, 17], or by using symmetry reductions [9]. These approaches however have worst case running time exponential in the number of processes in a system.

**Our Contribution** In this paper we follow the first approach to prevention of deadlock. We present two sufficient conditions for assuring deadlock freedom and describe efficient algorithms for checking these conditions. The first condition is suitable for concurrent programs consisting of similar processes, while the second applies to arbitrary concurrent programs. Essentially, the first condition is a modification of the condition presented in [2] and studies the “wait-for-graph” for presence of supercycles (supercycles represent cyclical waiting patterns). The check for the condition presented in [2] has run-time complexity, which is exponential in the branching degree of state-nodes in the “wait-for-graph”. In this paper, we present a condition that can be verified in polynomial time in all the input parameters: the number of processes, the size of a single process, and the branching degree of state-nodes. Moreover, the space complexity of the check is polynomial in the size of a single process. Therefore, this condition can be efficiently checked even on very large concurrent programs. We note that the condition in [2] is formulated for systems of similar (isomorphic) processes. However, it is easy to see that the restriction to similar processes does not play any role in the proof of correctness, and thus can be lifted.

The second condition is more complex, and also more discriminating. This condition is based on studying the global wait-for graph that presents the wait-for relations between all local states. We note that, unlike the previous work on deadlock detection, the graph contains only local states, and thus its size is always polynomial in the number of processes and in the size of a single process. We prove that deadlock results in a cyclic waiting condition in the graph (aka “supercycle”). However, not all supercycles correspond to reachable deadlocked states. A supercycle can result from an inconsistent or unreachable state, and thus presence of supercycles alone is not sufficient to declare a deadlock. We present an efficient algorithm for detection of supercycles in the global wait-for graph and for checking the discovered supercycles for inconsistency and reachability. For each supercycle discovered in the graph, we present an efficient algorithm for detection of inconsistencies and unreachability.

To the best of our knowledge, this is the first work that describes sufficient and polynomially checkable conditions for deadlock avoidance in systems of similar and dissimilar processes. We have implemented our pairwise representation using the XSB logic programming system [18]. This implementation provides a platform for implementing the algorithms in this paper.

## 2 Technical Preliminaries

### 2.1 Model of concurrent computation

We consider finite-state concurrent programs of the form  $P = P_1 \parallel \dots \parallel P_K$  that consist of a finite number  $n$  of fixed sequential processes  $P_1, \dots, P_K$  running in parallel. Each  $P_i$  is a *synchronization skeleton* [8], that is, a directed multigraph where each node is a (local) state of  $P_i$  (also called an *i-state* and is labeled by a unique name  $(s_i)$ ), and where each arc is labeled with a guarded command [7]  $B_i \rightarrow A_i$  consisting of a guard  $B_i$  and corresponding action  $A_i$ . With each  $P_i$  we associate a set  $\mathcal{AP}_i$  of *atomic propositions*, and a mapping  $V_i$  from local states of  $P_i$  to subsets of  $\mathcal{AP}_i$ :  $V_i(s_i)$  is the set of atomic propositions that are true in  $s_i$ . As  $P_i$  executes transitions and changes its local state, the atomic propositions in  $\mathcal{AP}_i$  are updated. Different local states of  $P_i$  have different truth assignments:  $V_i(s_i) \neq V_i(t_i)$  for  $s_i \neq t_i$ .

Atomic propositions are not shared:  $\mathcal{AP}_i \cap \mathcal{AP}_j = \emptyset$  when  $i \neq j$ . Other processes can read (via guards) but not update the atomic propositions in  $\mathcal{AP}_i$ . There is also a set of shared variables  $x_1, \dots, x_m$ , which can be read and written by every process. These are updated by the action  $A_i$ . A *global state* is a tuple of the form  $(s_1, \dots, s_K, v_1, \dots, v_m)$  where  $s_i$  is the current local state of  $P_i$  and  $v_1, \dots, v_m$  is a list giving the current values of  $x_1, \dots, x_m$ , respectively. A guard  $B_i$  is a predicate on global states, and so can reference any atomic proposition and any shared variable. An action  $A_i$  is a parallel assignment statement that updates the shared variables. We write just  $A_i$  for  $true \rightarrow A_i$  and just  $B_i$  for  $B_i \rightarrow skip$ , where *skip* is the empty assignment.

We model parallelism as usual by the nondeterministic interleaving of the “atomic” transitions of the individual processes  $P_i$ . Let  $s = (s_1, \dots, s_i, \dots, s_K, v_1, \dots, v_m)$  be the current global state, and let  $P_i$  contain an arc from node  $s_i$  to  $s'_i$  labeled with  $B_i \rightarrow A_i$ . We write such an arc as the tuple  $(s_i, B_i \rightarrow A_i, s'_i)$ , and call it a  $P_i$ -move from  $s_i$  to  $s'_i$ . We use just *move* when  $P_i$  is specified by the context. If  $B_i$  holds in  $s$ , then a permissible next state is  $s' = (s_1, \dots, s'_i, \dots, s_K, v'_1, \dots, v'_m)$  where  $v'_1, \dots, v'_m$  are the new values for the shared variables resulting from action  $A_i$ . Thus, at each step of the computation, a process with an enabled arc is nondeterministically selected to be executed next. The *transition relation*  $R$  is the set of all such  $(s, s')$ . The arc from node  $s_i$  to  $s'_i$  is *enabled* in state  $s$ . An arc that is not enabled is *blocked*.

Let  $S^0$  be a given set of initial states in which computations of  $P$  can start. A *computation path* is a sequence of states whose first state is in  $S^0$  and where each successive pair of states is related by  $R$ . A state is *reachable* iff it lies on some computation path. Let  $S$  be the set of all reachable global states of  $P$ , and redefine  $R$  to restrict it to  $S \times S$ , i.e., to reachable states. Then,  $M = (S^0, S, R)$  is the *global state transition diagram* (GSTD) of  $P$ .

## 2.2 Pairwise normal form

We will restrict our attention to concurrent programs that are written in a certain syntactic form, as follows. Let  $\oplus, \otimes$  be binary infix operators. A *general guarded command* [2] is either a guarded command as given in Section 2.1 above, or has the form  $G_1 \oplus G_2$  or  $G_1 \otimes G_2$ , where  $G_1, G_2$  are general guarded commands. Roughly, the operational semantics of  $G_1 \oplus G_2$  is that either  $G_1$  or  $G_2$ , but not both, can be executed, and the operational semantics of  $G_1 \otimes G_2$  is that both  $G_1$  or  $G_2$  must be executed, that is, the guards of both  $G_1$  and  $G_2$  must hold at the same time, and the bodies of  $G_1$  and  $G_2$  must be executed simultaneously, as a single parallel assignment statement. For the semantics of  $G_1 \otimes G_2$  to be well-defined, there must be no conflicting assignments to shared variables in  $G_1$  and  $G_2$ . This will always be the case for the programs we consider. We refer the reader to [2] for a comprehensive presentation of general guarded commands.

A concurrent program  $P = P_1 \parallel \dots \parallel P_K$  is in *pairwise normal form* iff the following four conditions all hold: (1) every move  $a_i$  of every process  $P_i$  has the form  $a_i = (s_i, \bigoplus_{j \in I(i)} \bigoplus_{\ell \in \{1, \dots, n_j\}} B_{i,\ell}^j \rightarrow A_{i,\ell}^j, t_i)$ , where  $B_{i,\ell}^j \rightarrow A_{i,\ell}^j$  is a guarded command,  $I$  is an irreflexive symmetric relation over  $\{1 \dots K\}$  that defines a “interconnection” (or “neighbors”) relation amongst processes, and  $I(i) = \{j \mid (i, j) \in I\}$ , (2) variables are shared in a pairwise manner, i.e., for each  $(i, j) \in I$ , there is some set  $\mathcal{SH}_{ij}$  of shared variables that are the only variables that can be read and written by both  $P_i$  and  $P_j$ , (3)  $B_{i,\ell}^j$  can reference only variables in  $\mathcal{SH}_{ij}$  and atomic propositions in  $\mathcal{AP}_j$ , and (4)  $A_{i,\ell}^j$  can update only variables in  $\mathcal{SH}_{ij}$ .

For each neighbor  $P_j$  of  $P_i$ ,  $\bigoplus_{\ell \in [1:n]} B_{i,\ell}^j \rightarrow A_{i,\ell}^j$  specifies  $n$  alternatives  $B_{i,\ell}^j \rightarrow A_{i,\ell}^j$ ,  $1 \leq \ell \leq n$  for the interaction between  $P_i$  and  $P_j$  as  $P_i$  transitions from  $s_i$  to  $t_i$ .  $P_i$  must execute such an interaction

with each of its neighbors in order to transition from  $s_i$  to  $t_i$ . We emphasize that  $I$  is not necessarily the set of all pairs, i.e., there can be processes that do not directly interact by reading each others atomic propositions or reading/writing pairwise shared variables. We do not assume, unless otherwise stated, that processes are isomorphic, or *similar* (we define process similarity later in this section).

We will usually use a superscript  $I$  to indicate the relation  $I$ , e.g., process  $P_i^I$ , and  $P_i^I$ -move  $a_i^I$ . We define  $a_i^I.start = s_i$ ,  $a_i^I.guard_j = \bigvee_{\ell \in \{1, \dots, n_j\}} B_{i,\ell}^j$ , and  $a_i^I.guard = \bigwedge_{j \in I(i)} a_i^I.guard_j$ . If  $P^I = P_1^I \parallel \dots \parallel P_K^I$  is a concurrent program with interconnection relation  $I$ , then we call  $P^I$  an  $I$ -system.

In pairwise normal form, the synchronization code for  $P_i^I$  with one of its neighbors  $P_j^I$  (i.e.,  $\bigoplus_{\ell \in \{1, \dots, n_j\}} B_{i,\ell}^j \rightarrow A_{i,\ell}^j$ ) is expressed separately from the synchronization code for  $P_i^I$  with another neighbor  $P_k^I$  (i.e.,  $\bigoplus_{\ell \in \{1, \dots, n_k\}} B_{i,\ell}^k \rightarrow A_{i,\ell}^k$ ). We can exploit this property to define “subsystems” of an  $I$ -system  $P$  as follows. Let  $J \subseteq I$  and  $range(J) = \{i \mid \exists j : (i, j) \in J\}$ . If  $a_i^I$  is a move of  $P_i^I$  then define  $a_i^J = (s_i, \bigoplus_{j \in J(i)} \bigoplus_{\ell \in \{1, \dots, n_j\}} B_{i,\ell}^j \rightarrow A_{i,\ell}^j, t_i)$ . Then the  $J$ -system  $P^J$  is  $P_{j_1}^J \parallel \dots \parallel P_{j_n}^J$  where  $\{j_1, \dots, j_n\} = range(J)$  and  $P_j^J$  consists of the moves  $\{a_i^J \mid a_i^I \text{ is a move of } P_j^I\}$ . Intuitively, a  $J$ -system consists of the processes in  $range(J)$ , where each process contains only the synchronization code needed for its  $J$ -neighbors, rather than its  $I$ -neighbors. If  $J = \{\{i, j\}\}$  for some  $i, j$  then  $P_J$  is a *pair-system*, and if  $J = \{\{i, j\}, \{j, k\}\}$  for some  $i, j, k$  then  $P_J$  is a *triple-system*. For  $J \subseteq I$ ,  $M_J = (S_J^0, S_J, R_J)$  is the GSTD of  $P^J$  as defined in Section 2.1, and a global state of  $P^J$  is a  $J$ -state. If  $J = \{\{i, j\}\}$ , then we write  $M_{ij} = (S_{ij}^0, S_{ij}, R_{ij})$  instead of  $M_J = (S_J^0, S_J, R_J)$ .

Also, if  $s_J$  is a  $J$ -state, and  $J' \subseteq J$ , then  $s \upharpoonright J'$  is the  $J'$ -state that agrees with  $s$  on the local state of all  $P_j \in range(J')$  and the value of all variables  $x_{ij} \in \mathcal{H}_{ij}$  such that  $i, j \in range(J')$ , i.e, the projection of  $s$  onto the processes in  $J'$ . If  $J' = \{\{i, j\}\}$  then we write  $s \upharpoonright J$  as  $s \upharpoonright ij$ . Also,  $s \upharpoonright i$  is the local state of  $P_i$  in  $s$ . Let  $reachable(P_i) = \{s_i \mid \exists j \in I(i), s_{ij} \in S_{ij} : s_{ij} \upharpoonright i = s_i\}$ , that is,  $reachable(P_i)$  is the set of local states of  $P_i$  that are reachable in some pair-system involving  $P_i$ .

Two processes  $P_i$  and  $P_j$  are *similar* if they are isomorphic to each other up to the change of indices [2]. A concurrent program  $P = P_1 \parallel \dots \parallel P_K$  consists of *similar processes* if for each  $1 \leq i, j \leq K$ , we have that  $P_i$  and  $P_j$  are similar.

[2, 1, 3] give, in pairwise normal form, solutions to many well-known problems, such as dining philosophers, drinking philosophers, mutual exclusion,  $k$ -out-of- $n$  mutual exclusion, two-phase commit, and replicated data servers. Attie [4] shows that any finite-state concurrent program can be rewritten (up to strong bisimulation) in pairwise normal form. The restriction to pairwise normal form enables us to mechanically verify certain correctness properties very efficiently. For example, safety properties that can be expressed over pairs of processes can be verified in time  $O(K^2 N^2)$ , where  $K$  is the number of processes and  $N$  is the size of the largest process. Exhaustive state-space enumeration would of course require  $O(N^K)$  time. In this paper, we exploit pairwise normal form to devise sufficient conditions for deadlock-freedom that can be checked in time polynomial in both  $N$  and  $K$ .

### 2.3 The Wait-For-Graph

The wait-for-graph for an  $I$ -state  $s$  gives all of the blocking relationships in  $s$ .

**Definition 2.1** (Wait-For-Graph  $W_I(s)$ ) *Let  $s$  be an arbitrary  $I$ -state. The wait-for-graph  $W_I(s)$  of  $s$  is a directed bipartite AND-OR graph, where*

1. The AND nodes of  $W_I(s)$  (also called local-state nodes) are the  $i$ -states  $\{s_i \mid s_i = s \upharpoonright i, i \in \{1 \dots K\}\}^1$ ;
2. The OR-nodes of  $W_I(s)$  (also called move nodes) are the moves  $\{a_i^I \mid i \in \{1 \dots K\}$  and  $a_i^I$  is a move of  $P_i^I$  and  $s \upharpoonright i = a_i^I.start\}$
3. There is an edge from  $s_i$  to every node of the form  $a_i^I$  in  $W_I(s)$ ;
4. There is an edge from  $a_i^I$  to  $s_j$  in  $W_I(s)$  if and only if  $\{i, j\} \in I$  and  $a_i^I \in W_I(s)$  and  $s \upharpoonright ij(a_i^I.guard_j) = false$ .

The AND-nodes are the local states  $s_i = s \upharpoonright i$  of all processes when the global state is  $s$ , and the OR-nodes are the moves  $a_i^I$  such that local control in  $P_i^I$  is currently at the start state of  $a_i^I$ , i.e., all the moves that are candidates for execution. There is an edge from  $s_i$  to each move of the form  $a_i^I$ . Nodes  $s_i$  are AND nodes since  $P_i^I$  is blocked if and only if *all* of its possible moves are blocked. There is an edge from  $a_i^I$  to  $s_j$  iff  $a_i^I$  is blocked by  $P_j^I$ :  $a_i^I$  can be executed in  $s$  only if  $s \upharpoonright ij(a_i^I.guard_j) = true$  for all  $j \in I(i)$ ; if there is some  $j$  in  $I(i)$  such that  $s \upharpoonright ij(a_i^I.guard_j) = false$ , then  $a_i^I$  cannot be executed in state  $s$ . The nodes labeled with moves are OR nodes, since  $a_i^I$  is blocked iff *some* neighbor  $P_j^I$  of  $P_i^I$  blocks  $a_i^I$ . Note, however, that we cannot say that  $P_i^I$  itself is blocked by  $P_j^I$ , since there could be another move  $b_i^I$  in  $P_i^I$  such that  $s \upharpoonright ij(b_i^I.guard_j) = true$ , i.e.,  $b_i^I$  is not blocked by  $P_j^I$  (in state  $s$ ), so  $P_i^I$  can progress in state  $s$  by executing  $b_i^I$ .

In the sequel, we use  $s_i \longrightarrow a_i^I \in W_I(s)$  to denote the existence of an edge from  $s_i$  to  $a_i^I$  in  $W_I(s)$ , and  $a_i^I \longrightarrow s_j \in W_I(s)$  to denote the existence of an edge from  $a_i^I$  to  $s_j$  in  $W_I(s)$ . We also abbreviate  $((s_i \longrightarrow a_i^I \in W(s)) \wedge (a_i^I \longrightarrow s_j \in W(s)))$  with  $s_i \longrightarrow a_i^I \longrightarrow s_j \in W(s)$ , and similarly for longer “wait-chains.” For  $J \subseteq I$  and  $J$ -state  $s_J$  we define  $W_J(s_J)$  by replacing  $I$  by  $J$  and  $\{1 \dots K\}$  by  $range(J)$  in the above definition.

## 2.4 Establishing Deadlock-freedom: Supercycles

Deadlock is characterized by the presence in the wait-for-graph of a graph-theoretic construct called a *supercycle* [2]:

**Definition 2.2 (Supercycle)** *SC is a supercycle in  $W_I(s)$  if and only if all of the following hold:*

1. *SC is nonempty,*
2. *if  $s_i \in SC$  then  $\forall a_i^I : a_i^I \in W_I(s)$  implies  $s_i \longrightarrow a_i^I \in SC$ , and*
3. *if  $a_i^I \in SC$  then  $\exists s_j : a_i^I \longrightarrow s_j \in W_I(s)$  and  $a_i^I \longrightarrow s_j \in SC$ .*

Note that  $SC$  is a subgraph of  $W_I(s)$ . If an  $i$ -state  $s_i$  is in a supercycle  $SC$ , then every move of  $P_i^I$  that starts in  $s_i$  is also in  $SC$  and is blocked by some other  $I$ -process  $P_j^I$  which has a  $j$ -state  $s_j$  in  $SC$  (note that a process has at most one local state in  $SC$ , and we say that the process itself is in  $SC$ ). It follows that no  $I$ -process in  $SC$  can execute any of its moves, and that this situation persists forever.

In Figure 1 we give an example of a wait-for-graph for a three process system. And-nodes (processes) are shown as  $\bullet$ , and or-nodes (moves) are shown as  $\circ$ . Each process  $P_i, i \in \{1, 2, 3\}$  has two moves  $a_i$

<sup>1</sup>In [2] state nodes are denoted by processes  $P_i$  and not by local states, since they consider wait-for-graphs for each state of the system separately; in this paper, we study wait-for-graphs that encompass all blocking conditions for all local nodes of all processes together; hence we need to distinguish between different local state-nodes of the same process.

and  $b_i$ . Since every move has at least one outgoing edge, i.e., is blocked by at least one process, the figure also an example of a supercycle. In fact, several edges can be removed and still leave a supercycle (for example,  $a_3 \rightarrow P_1$ ,  $b_3 \rightarrow P_2$ ,  $a_2 \rightarrow P_1$  can all be removed). Thus, the figure contains several subgraphs that are also supercycles.




Figure 1: Example of a wait-for-graph.

From [2], we have that the absence of supercycles in the wait-for-graph of a state implies that there is at least one enabled move in that state:

**Proposition 2.3 ([2])** *If  $W_I(s)$  is supercycle-free, then some move  $a_i^I$  has no outgoing edges in  $W_I(s)$ , and so can be executed in state  $s$ .*

We say that  $s$  is *supercycle-free* iff  $W_I(s)$  does not contain a supercycle. We assume that all initial states of the  $I$ -system are supercycle free. That is, we do not allow initial states that contain deadlocks.

### 3 Improving the deadlock freedom condition from [2]

#### 3.1 The wait-for-graph condition

Consider the following condition:

$$\begin{aligned} &\text{For every reachable } I\text{-state } t \text{ in } M_I \\ &\text{such that } s \xrightarrow{k} t \in R_I \text{ for some reachable } I\text{-state } s, \\ &\quad \neg \exists a_j^I : (a_j^I \rightarrow P_k^I \in W_I(t)) \text{ or} \\ &\quad \exists a_k^I \in W_I(t) : (\bigwedge_{\ell \in \{1 \dots K\}} (a_k^I \rightarrow P_\ell^I \notin W_I(t))). \end{aligned} \tag{a}$$

This condition implies that, after  $P_k^I$  executes a transition, either  $P_k^I$  blocks no move of another process, or  $P_k^I$  itself has an enabled move. Thus  $P_k^I$  cannot be in a supercycle. Hence, this transition of  $P_k^I$  could not have *created* a supercycle; any supercycle present after the transition must also have been present before the transition. Since initial states are supercycle-free, we conclude, by induction on computation path length, that every reachable  $I$ -state is supercycle-free.

Let  $t_k.\text{moves}$  denote the set of moves  $\{a_k^I \mid a_k^I \in P_k^I \text{ and } a_k^I.\text{start} = t_k\}$ . It is proved in [2] that it is enough to check condition (a) for all  $J$ -systems for  $J \in \mathcal{J}$ , where  $\mathcal{J}$  is the set of all interconnection relations of the form  $\{\{j, k\}, \{k, \ell_1\}, \{k, \ell_2\}, \dots, \{k, \ell_n\}\}$ , where  $n = |(t|k).\text{moves}|$ ,

$1 \leq j, k, \ell_1 \dots, \ell_n \leq K, k \notin \{j, k, \ell_1 \dots, \ell_n\}$ . That is, we check:

$$\begin{aligned}
& \text{For every } t_k \in \text{reachable}(P_k), \\
& \text{for all } J \in \mathcal{J} \text{ in which there exists a reachable } J\text{-state } t_J \\
& \text{such that } s_J \xrightarrow{k} t_J \text{ for some } J\text{-reachable state } s_J \text{ and } t_J \upharpoonright k = t_k, \\
& \quad \neg \exists a_j^J : (a_j^J \longrightarrow P_k^J \in W_J(t_J)) \text{ or} \\
& \quad \exists a_k^J \in W_J(t_J) : (\bigwedge_{\ell \in \{\ell_1, \dots, \ell_n\}} (a_k^J \longrightarrow P_\ell^J \notin W_J(t_J))).
\end{aligned} \tag{b}$$

This condition implies an algorithm that checks all possible subsystems  $J$  of the form  $\{\{j, k\}, \{k, \ell_1\}, \dots, \{k, \ell_n\}\}$ . The algorithm is exponential in  $n$ , thus for large  $n$  constructing the subsystems is infeasible.

In the condition (b), the variables that appear under the disjunction over the moves of  $P_k^J$  in  $t_k$  ( $\exists a_k^J \in W_J(t_J)$ ) do not involve  $a_j^J$ , and so we can move the disjunction over the moves  $a_k^J$  of  $P_k^J$  before the conjunction over the moves  $a_j^J$  of  $P_j^J$  ( $\neg \exists a_j^J$ ). We get the following equivalent condition:

$$\begin{aligned}
& \text{For every } t_k \in \text{reachable}(P_k), \\
& \text{for all } J \in \mathcal{J} \text{ in which there exists a reachable } J\text{-state } t_J \\
& \text{such that } s_J \xrightarrow{k} t_J \text{ for some } J\text{-reachable state } s_J \text{ and } t_J \upharpoonright k = t_k, \\
& \quad \exists a_k^J \in W_J(t_J) : \\
& \quad \neg \exists a_j^J : (a_j^J \longrightarrow P_k^J \in W_J(t_J)) \text{ or } (\bigwedge_{\ell \in \{\ell_1, \dots, \ell_n\}} (a_k^J \longrightarrow P_\ell^J \notin W_J(t_J))).
\end{aligned}$$

Let  $J_i = \{\{j, k\}, \{k, \ell_i\}\} \subseteq J$ , for  $1 \leq i \leq n$ . Then, for each move  $a_k^J, \bigwedge_{\ell \in \{\ell_1, \dots, \ell_n\}} (a_k^J \rightarrow P_\ell^J \notin W_J(t_J))$  holds iff

$$\bigwedge_{J_i: 1 \leq i \leq n} (a_k^{J_i} \rightarrow P_{\ell_i}^{J_i} \notin W_{J_i}(t_J \upharpoonright J_i)), \tag{1}$$

The last equation follows from wait-for-graph projection [2, Proposition 6.5.4.1].

Equation 1 is checked with respect to all systems of three processes, for all reachable states of the  $J$ -system. In order to avoid constructing the  $J$ -system, we check the following condition, which requires only construction of the  $J_i$ -systems.

$$\begin{aligned}
& \text{For every } t_k \in \text{reachable}(P_k), \\
& \text{there exists a move } a_k^J \in t_k.\text{moves} \\
& \text{such that for all relations } J_i = \{\{j, k\}, \{k, \ell_i\}\} \subseteq J \\
& \text{in which there exists a reachable state } t_{J_i} \\
& \text{such that } s_{J_i} \xrightarrow{k} t_{J_i} \text{ for some } J_i\text{-reachable state } s_{J_i} \text{ and } t_{J_i} \upharpoonright k = t_k, \\
& \quad \text{we have } \neg \exists a_j^{J_i} : (a_j^{J_i} \longrightarrow P_k^{J_i} \in W_{J_i}(t_{J_i})) \text{ or} \\
& \quad (a_k^{J_i} \longrightarrow P_{\ell_i}^{J_i} \notin W_{J_i}(t_{J_i})).
\end{aligned} \tag{c}$$

In other words, condition (c) holds if either  $P_k$  blocks no move of another process or there exists a move of  $P_k$  that is not blocked in any of the triple systems  $J_i$ .

**Theorem 3.1** *If condition (c) holds, then (1) condition (b) holds, and (2) the I-system is deadlock-free.*

**Proof:** Condition (b) implies deadlock-freedom [2]. Thus, it suffices to prove that the condition (c) implies the condition (b). We prove the contrapositive. Assume that  $M_J$  does not satisfy condition (b) in a state  $t_k^J$ . Then, there exists a move  $a_j^J$  of  $P_j$  such that  $a_j^J \rightarrow P_k^J \in W_J(t_J)$ , and for each move  $a_k^J \in (t_k^J \upharpoonright k).\text{moves}$ , there exists a process  $P_\ell^J$  such that  $a_k^J \rightarrow P_\ell^J \in W_J$ . Let  $\{a_{k,1}^J, \dots, a_{k,n}^J\} =$



$(t_k^J \upharpoonright k).moves$ . For a move  $a_{k,i}^J$ , let  $P_{l_i}$  be the process that blocks it in  $t_k^J$ , for  $1 \leq i \leq n$ . Now we show that condition (c) is violated. Indeed, for each interprocess relation  $J_i = \{\{j, k\}, \{k, l_i\}\}$ , for  $1 \leq i \leq n$ , we have that  $P_k^{J_i}$  blocks the move  $a_j^{J_i}$  of the process  $P_j^{J_i}$ . In addition, for each move  $a_{k,i}$  of  $P_k$  there exists a relation  $J_i = \{\{j, k\}, \{k, l_i\}\}$  in which  $a_{k,i}$  is blocked.  $\square$

Intuitively, checking condition (c) involves constructing all triples of processes with  $P_k$  being the middle process. Since the size of a triple system is polynomial in the size of a single process, and the number of triples is polynomial in the number of processes in the system, the check is polynomial in all parameters. Now we describe the algorithm formally and analyze its time and space complexity. For a process  $P_k$ , we determine the set  $S_k$  of reachable local states of  $P_k$ . For every  $t_k \in S_k$ , we determine the set  $t_k.moves$  of outgoing moves of  $P_k$  from  $t_k$ . The  $J_i$ -systems in the wait-for-graph assumption are of the form  $J_i = \{\{j, k\}, \{k, l_i\}\}$ , where  $k \neq j$  and  $k \neq l_i$ . We define the set of all  $J$ -systems under consideration as

$$\mathcal{J}(t_k) = \{J_i : J_i = \{\{j, k\}, \{k, l_i\}\}, k \neq j, l_i \text{ and } l_i \in \{l_1, \dots, l_n\}\},$$

where  $n = |t_k.moves|$ .

For every  $J_i \in \mathcal{J}$  we generate the state-transition diagram  $M_{J_i}$  and the wait-for-graph  $W_{J_i}(t_k^{J_i})$ . For each  $a_k \in t_k.moves$ , we evaluate

$$\bigwedge a_j^{J_i} : (a_j^{J_i} \rightarrow P_k^{J_i} \notin W_{J_i}(t_k^{J_i})) \vee (a_k^{J_i} \rightarrow P_{l_i}^{J_i} \notin W_{J_i}(t_k^{J_i})) \quad (2)$$

for all  $J_i \in \mathcal{J}(t_k)$ , where  $a_k^{J_i} \upharpoonright k = a_k$ .

If for all  $t_k$  there exists  $a_k \in t_k.moves$  for which Equation 2 holds for all  $J_i \in \mathcal{J}$ , then we conclude that the system is deadlock-free. We formalize the procedure given above in Figure 2, and compute its time complexity. The procedure clearly terminates, since the range of all loop variables is finite. Furthermore, upon termination, condition (c) holds if and only if  $WGflag$  is set to “true.” By definition of  $M_{J_i}$ , its size is bounded by  $O(N^3)$ , where  $N$  is the size of a single process.

We consider each step of the procedure and compute its worst-case time complexity. The complexity of step 2.3.2.2 is  $O(|M_{J_i}|)$ , which is  $O(N^3)$ . The step 2.3.2.2 is performed  $|\mathcal{J}| = O(K^3)$  times. The step 2.3.2 is performed  $|t_k.moves| = n$  times. Thus, the total runtime complexity is  $O(K^3 N^3 n)$ , which is polynomial in the number of processes, the size of a single process, and the branching degree of a process. The space complexity is bounded by the size of a triple system, which is  $O(N^3)$ .

**Example 3.2** [Deadlock detection in the general resource allocation problem] In this example we apply our method to the general resource allocation problem presented in Lynch [16, Chapter 11]. The problem is defined by the *explicit resource specification* for  $n$  processes. An explicit resource specification  $\mathcal{R}$  for  $n$  processes consists of a universal finite set  $R$  of (unsharable) resources and sets  $R_i \subseteq R$  for all  $1 \leq i \leq n$ , where  $R_i$  is the set of resources that process  $P_i$  requires to execute. We describe here an instance of the resource allocation problem in which there is a potential deadlock and show how this deadlock can be detected by studying triples of processes. We assume that each process needs at least one resource in order to execute. We first consider a naive algorithm in which each process chooses the order of requests for resources non-deterministically. That is, if a process  $P_i$  needs resources  $\{1, \dots, k\}$ , it non-deterministically acquires resource  $1 \leq r_1 \leq k$ , then a resource  $r_2 \in \{1, \dots, k\} \setminus r_1$ , etc. After the last resource has been acquired,  $P_i$  executes. Clearly, if a resource  $r$  is already allocated to another

0.  $WGflag := true;$
1.  $S_k := reachable(P_k);$
2. for all  $t_k \in S_k$ 
  - 2.1  $n := |t_k.moves|;$
  - 2.2  $\mathcal{J}(t_k) := \{J_i \mid J = \{\{j, k\}, \{k, l_i\}\}\};$
  - 2.3 for all  $a_k \in t_k.moves$ 
    - 2.3.1 reset the value of Equation 2 to true;
    - 2.3.2 for all  $J_i$  in  $\mathcal{J}(t_k)$ 
      - 2.3.2.1 generate  $M_{J_i};$
      - 2.3.2.2 evaluate Equation 2; if false, break;
  - 2.4 if for all  $a_k$  the value of Equation 2 is false, set  $WGflag$  to false; return;

Figure 2: Procedure that checks the wait-for-graph assumption.

process,  $P_i$  cannot acquire it. If at some state in the resources allocation all remaining resources are allocated to other processes,  $P_i$  cannot proceed.

We show that in the case where  $R_i = R$  for all  $1 \leq i \leq n$  (every process needs all resources) and  $|r| \leq n$ , the system has deadlocked states. The skeleton of a single process  $P_i$  is given in Figure 3. The states in which  $P_i$  requires resources are labeled with sets of resources already acquired by  $P_i$ . For sets  $S_1, S_2 \subseteq R$  of resources, there is a move from the state labeled with  $S_1$  to the state labeled with  $S_2$  iff  $S_2 = S_1 \cup \{r\}$  for some resource  $r$ , and the move is guarded by  $\otimes_{j \neq i} x_{ijr}^i = i$ . There is a move  $S \rightarrow C_i$  iff  $S = R$  (that is, a process can move to the state where it executes only after it acquired all its resources).

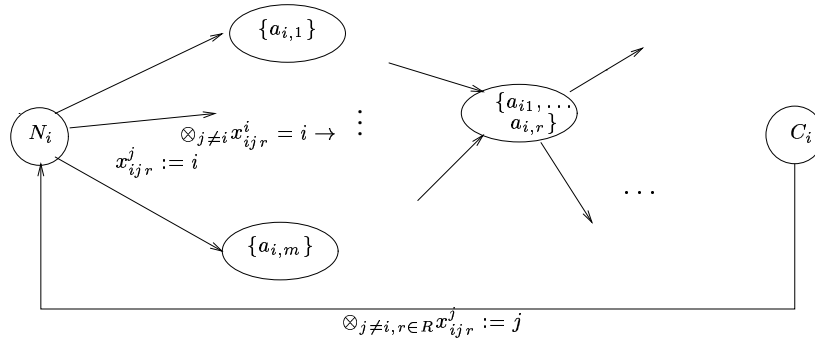


Figure 3: The skeleton of a process in the resource allocation problem

For  $1 \leq i, j \leq n$  and  $1 \leq r \leq |R|$ , the variables  $x_{ijr}^i$  are shared variables that are written by  $P_j$  and read by  $P_i$  only and determine the allocation of the resource  $r$ . A process  $P_i$  can acquire the resource  $r$  only if  $x_{ijr}^i = i$  for all  $1 \leq j \leq n$ . Upon acquiring the resource  $r$ , the process  $P_i$  sets  $x_{ijr}^j$  to  $i$ , for all  $1 \leq j \leq n$ . The state of  $P_i$  is determined by the subset of resources  $P_i$  currently holds. That is,

$P_i$  is in state  $\langle a_{i,1}, \dots, a_{i,l} \rangle$  iff  $P_i$  holds the resources  $\{r_1, \dots, r_l\}$ . The system has a single initial state  $[N_1 N_2 \dots N_n]$ . The process  $P_i$  executes in the state  $C_i$  for  $1 \leq i \leq n$ . Let  $|R| = m$ . For simplicity, we assume  $n = m$  (if  $n > m$  the deadlock follows from the deadlock for  $n = m$ ). The system reaches the deadlocked state  $[\langle a_{1,1} \rangle \langle a_{2,2} \rangle \dots \langle a_{m,m} \rangle]$  (i.e., the state in which the process  $P_i$  holds the resource  $r_i$  for  $1 \leq i \leq m$  via the following path.

$$\begin{aligned}
& [N_1 N_2 \dots N_m] \xrightarrow{1} \\
& [x_{1i1}^i = 1 \text{ for } i \in \{2, \dots, m\}, \langle a_{1,1} \rangle N_2 \dots N_m] \xrightarrow{2} \\
& [x_{1i1}^i = 1 \text{ for } i \in \{2, \dots, m\}, x_{2i2}^i = 2 \text{ for } i \in \{1, \dots, m\} \setminus \{2\}, \langle a_{1,1} \rangle \langle a_{2,2} \rangle N_3 \dots N_m] \xrightarrow{3} \dots \\
& \dots [x_{1i1}^i = 1 \text{ for } i \in \{2, \dots, m\}, \dots x_{mrm}^r = r \text{ for } r \in \{1, \dots, m-1\} \langle a_{1,1} \rangle \langle a_{2,2} \rangle \dots \langle a_{m,m} \rangle]
\end{aligned}$$

It is easy to see that the last state is a deadlock state, i.e., has no enabled outgoing moves. Indeed, every process holds exactly one resource, and thus no process can continue the allocation and execute. This scenario arises when the process  $P_1$  requests the first resource, the process  $P_2$  requests the second resource, etc. until the process  $P_m$  requests the  $m$ -th resource.

The deadlock can be detected by checking condition (c) in the  $J_{ik}$ -systems  $J_{ik} = \{\{i, m\}, \{m, k\}\}$ . In the state  $[\langle a_{1,1} \rangle \langle a_{2,2} \rangle \dots \langle a_{m,m} \rangle]$  of the  $J_{ik}$ -system the process  $P_m$  blocks the move  $\langle a_{i,i} \rangle \rightarrow \langle a_{i,i}, a_{i,m} \rangle$  of  $P_i$  (that is, the move that tries to allocate the resource  $m$  to  $P_i$ ), and the process  $P_k$  blocks the move  $\langle a_{m,m} \rangle \rightarrow \langle a_{m,m}, a_{m,k} \rangle$  of  $P_m$  by setting the variable  $x_{kmk}^m$  to  $k$  (that is, holding the resource  $k$ ). Thus, for each move  $\langle a_{m,m} \rangle \rightarrow \langle a_{m,m}, a_{m,k} \rangle$  there exists a triple-system  $J_{ik}$  in which this move is blocked, and therefore condition (c) fails. We note that the number of processes and the number of resources can be arbitrarily large. Indeed, checking condition (c) requires only constructing triples of processes, and the number of triples is polynomial in the total number of processes in the system.

Now consider the hierarchical resource allocation presented in Lynch [16, Chapter 11]. In this case, there is a global hierarchy between processes, and the resource is acquired to the process with the highest priority that requests it. The system is deadlock-free. However, condition (c) fails, giving a false deadlock indication. The reason for its failure is existence of waiting chains of length three in the system, despite the fact that cyclical waiting pattern never occurs. In the next section we present a more complex (and more discriminating) algorithm that shows deadlock freedom of hierarchical resource allocation.

We study one more example of deadlock detection, and then we modify the example so that, while condition (c) fails, the system is in fact deadlock-free. In Section 4 we show a more discriminating test, which shows deadlock-freedom of the example.

**Example 3.3** [Application of our algorithm to detect possible deadlock] Consider the system of  $K$  similar processes with a single initial state  $[N_1 \dots N_K]$ . The skeleton of a single process  $P_i$  is given in Figure 4. The processes attempt to enter one of the two critical sections  $C$  and  $D$ .

For  $i, j \in \{1 \dots K\}, i \neq j$ , the variables  $x_{ij}^i, y_{ij}^i$  are shared variables that are written by  $P_j$  and read by  $P_i$  only. The shared variable  $x_{ij}^i$  controls the move of  $P_i$  from  $T_i$  to  $C_i$ , which can be made only when  $x_{ij}^i = i$ . Likewise,  $y_{ij}^i$  controls the move of  $P_i$  from  $T_i$  to  $D_i$ , which can be made only when  $y_{ij}^i = i$ . When process  $P_j$  moves from the state  $N_j$  to the state  $T_j$ , it non-deterministically chooses how to update the variables  $x_{ij}^i, y_{ij}^i$ , for  $i \in \{1 \dots K\} \setminus \{j\}$ . In the state  $T_j$ , if both moves of  $P_j$  are enabled, it chooses the next move non-deterministically. For  $K = 4$ , the system can reach a deadlocked state  $[T_1, T_2, T_3, T_4]$  via the following path.




Figure 4: The skeleton of a process in the deadlock example for four processes.

$$\begin{aligned}
& [N_1 N_2 N_3 N_4] \xrightarrow{1} \\
& [x_{1i}^i = 1 \text{ for } i \in \{2, 3, 4\}, T_1 N_2 N_3 N_4] \xrightarrow{2} \\
& [x_{1i}^i = 1 \text{ for } i \in \{2, 3, 4\}, y_{2j}^j = 2 \text{ for } j \in \{1, 3, 4\}, T_1 T_2 T_3 T_4] \xrightarrow{3} \\
& [x_{1i}^i = 1 \text{ for } i \in \{2, 3, 4\}, y_{2j}^j = 2 \text{ for } j \in \{1, 3, 4\}, x_{3k}^k = 3 \text{ for } k \in \{1, 2, 4\}, T_1 T_2 T_3 N_4] \xrightarrow{4} \\
& [x_{1i}^i = 1 \text{ for } i \in \{2, 3, 4\}, y_{2j}^j = 2 \text{ for } j \in \{1, 3, 4\}, x_{3k}^k = 3 \text{ for } k \in \{1, 2, 4\}, \\
& y_{4l}^l = 4 \text{ for } l \in \{1, 2, 3\}, T_1 T_2 T_3 T_4]
\end{aligned}$$

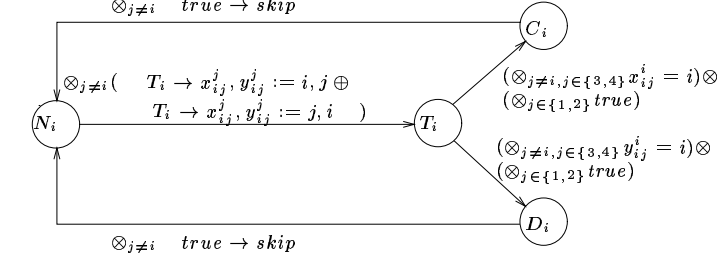
It is easy to see that the last state is a deadlock state, i.e., has no enabled outgoing moves. Indeed, the outgoing move  $T_1 \rightarrow C_1$  is blocked by the process  $P_3$ , and the outgoing move  $T_1 \rightarrow D_1$  is blocked by the processes  $P_2$  and  $P_4$ , and similarly for other processes. This can be detected by checking condition (c) in the  $J$ -system  $J = \{\{1, 4\}, \{4, 2\}\}$ . In the state  $[T_1 T_4 T_2]$  of this  $J$ -system the process  $P_4$  blocks the move  $T_1 \rightarrow D_1$  of  $P_1$ , and the process  $P_2$  can block the move  $T_4 \rightarrow C_4$  of  $P_4$  by setting the variable  $x_{24}^4$  to 2, or the move  $T_4 \rightarrow D_4$  by setting the variable  $y_{24}^4$  to 2. Thus, condition (c) fails for this  $J$ -system.

We note that  $P_2$  can block only one move of  $P_4$  at a time. However, since all processes are similar, in the  $K$ -process system  $P_2$  can block one move and another process can block another move of  $P_4$ , so that this deadlock scenario applies for any number  $K \geq 4$  of processes.

The requirement of processes similarity is therefore crucial. In Section 3.4 we describe an example of a system with non-similar processes and show that in this case we get a false deadlock indication from checking condition (c).


**Example 3.4** [False positive deadlock indication in a system with dissimilar processes] In this example we describe a situation in which the deadlock-freedom check described in Section 3 may be too conservative. Consider a modification of Example 3.3 for a system of four processes, in which the processes  $P_1$  and  $P_2$  cannot block the processes  $P_3$  and  $P_4$ , that is,  $P_i$  always assigns  $j$  to  $x_{ij}^j$  when  $i \in \{1, 2\}$  and  $j \in \{3, 4\}$ , rather than nondeterministically choosing between assigning  $i$  (thereby blocking  $P_j$ ) and assigning  $j$ . Condition (c) still fails in the state  $[T_1 T_2 T_3 T_4]$  with  $P_4$  being the process that made the last move. The system is not deadlocked, however, since the only process that can block  $P_4$  is  $P_3$ , and it can block only one move of  $P_4$  at the same time. Thus,  $P_4$  always has at least one enabled move. Figure 5 shows the skeletons of processes in this example. The system has a single initial state  $[N_1 N_2 N_3 N_4]$ . The processes  $P_1$  and  $P_2$  access the critical sections  $A$  and  $B$ , and the processes  $P_3$  and  $P_4$  access the critical sections  $C$  and  $D$ . The processes  $P_1$  and  $P_2$  can block each other, and the processes  $P_3$  and  $P_4$

can block each other and the processes  $P_1$  and  $P_2$ . Condition (c) fails in the state  $[T_1 T_2 T_3 T_4]$  with  $P_4$  being the process that made the last move. Indeed, in the system  $J = \{\{1, 4\}, \{4, 3\}\}$ , the process  $P_4$  blocks the process  $P_1$  and both moves of  $P_4$ , when viewed separately, can be blocked by  $P_3$ . However, the system is not deadlocked in the state  $[T_1 T_2 T_3 T_4]$  for all reachable combinations of values of the shared variables. The process  $P_3$  can block only one move of  $P_4$  at the same time, and the processes  $P_1$  and  $P_2$  cannot block  $P_4$ . Thus,  $P_4$  always has at least one enabled move.



$P_i$  for  $i \in \{3, 4\}$ .

The processes  $P_3$ , and  $P_4$  can block each other and the processes  $P_1$  and  $P_2$



$P_i$  for  $i \in \{1, 2\}$ .

The processes  $P_1$  and  $P_2$  can only block each other

Figure 5: Modified example for four dissimilar processes.

The test given in the next section, which is more complex but also more discriminating, verifies that this example is in fact deadlock-free.

## 4 More Complex and Discriminating Condition

We define a *global wait-for graph*  $\mathcal{W}$  which contains the union of all  $W_I(s)$ , for all reachable  $I$ -states  $s$ .

**Definition 4.1** ( $\mathcal{W}$ ) *The graph  $\mathcal{W}$  is as follows. The nodes of  $\mathcal{W}$  are*

1. *the states  $s_i$  such that  $i \in \{1 \dots K\}$  and  $s_i \in \text{reachable}(P_i)$ ;*
2. *the moves  $a_i$  such that  $i \in \{1 \dots K\}$ ,  $a_i$  is a move of  $P_i$ , and  $a_i.\text{start} = s_i$  for some node  $s_i$ ;*

*and the edges are:*

1. *an edge from  $s_i$  to every  $a_i$  such that  $a_i.\text{start} = s_i$ ;*
2. *for  $(i, j) \in I$  and  $a_i$  a move  $a_i$  of  $P_i$ , there is an edge from  $a_i$  to  $s_j$  iff  $\exists s_{ij} \in S_{ij} : s_{ij} \downarrow j = s_j \wedge s_{ij}(a_i^I.\text{guard}_j) = \text{false}$ .*

We can view  $\mathcal{W}$  as either a directed graph or as an AND-OR graph. When viewed as an AND-OR graph, the AND-nodes are the local states  $s_i$  of all processes (which we call local-state nodes) and the OR-nodes are the moves  $a_i$  (which we call move nodes). We shall use MSCC to abbreviate “maximal strongly connected component” in the sequel. The following proposition is an immediate consequence of the definitions.

**Proposition 4.2** *For every reachable I-state  $s$ ,  $W_I(s)$  is a subgraph of  $\mathcal{W}$ .*

**Proposition 4.3** *Let  $s$  be a reachable I-state, and assume that  $W_I(s)$  contains a supercycle  $SC$ . Then, there exists a subgraph  $SC'$  of  $SC$  such that:*

1.  $SC'$  is nontrivial, i.e., contains more than one node,
2.  $SC'$  is itself a supercycle, and
3. there exists a maximal strongly connected component  $\mathcal{W}'$  of  $\mathcal{W}$  such that  $SC'$  is a subgraph of  $\mathcal{W}'$ .

**Proof:** By Definition 2.2,  $SC$  is not acyclic, hence  $SC$  contains at least one (nontrivial) MSCC. Consider the acyclic component graph obtained by shrinking each MSCC down to a single node. This graph, being acyclic, contains at least one node with no outgoing edges. Let  $SC'$  be an MSCC of  $SC$  corresponding to such a node. Thus, no node in  $SC'$  has an outgoing edge to a node outside of  $SC'$ . Suppose  $SC'$  contains a single local-state node  $s_i$ . Then, by Definition 2.2 and our restriction (in Section 2.1) that a synchronization skeleton contains no “dead ends,”  $SC'$  must contain at least one node of the form  $a_i$ . On the other hand, suppose that  $SC'$  contains a single move node  $a_i$ . Since  $a_i$  is not enabled in global state  $s$ , it must be blocked by at least one local-state node  $s_j (= s \downarrow j)$ . Hence,  $s_j$  is a node in  $SC'$ . Thus, in either case,  $SC'$  contains at least two nodes. Thus clause (1) is established.

We showed above that no node in  $SC'$  has an outgoing edge to a node outside of  $SC'$ . Thus, every move-node in  $SC'$  is blocked only by local-state nodes in  $SC'$ . Hence, by clause (1) and Definition 2.2,  $SC'$  itself is a supercycle. Hence, clause (2) is established.

Suppose  $SC'$  does not occur within a single MSCC of  $\mathcal{W}$ . Then there exist two MSCC's  $\mathcal{W}_1, \mathcal{W}_2$  of  $\mathcal{W}$ , and two nodes  $v_1, v_2$  such that  $v_1$  occurs in both  $SC'$  and  $\mathcal{W}_1$  and  $v_2$  occurs in both  $SC'$  and  $\mathcal{W}_2$ . Since  $SC'$  is a strongly connected subgraph of  $\mathcal{W}$ ,  $v_1$  is reachable from  $v_2$  in  $\mathcal{W}$  and vice-versa. Hence  $\mathcal{W}_1$  and  $\mathcal{W}_2$  are not *maximal* strongly connected components, contrary to assumption. Hence clause (3) is established.  $\square$

**Proposition 4.4** *If  $\mathcal{W}$  is acyclic, then for all reachable I-states  $s$ ,  $W_I(s)$  is supercycle-free.*

**Proof:** We establish the contrapositive. Suppose  $W_I(s)$  contains a supercycle  $SC$  for some  $s \in S_I^r$ . Then by Proposition 4.2,  $SC$  is a subgraph of  $\mathcal{W}$ . But  $SC$  contains a cycle. Hence  $\mathcal{W}$  is not acyclic.  $\square$

We now present a test for supercycle-freeness. In the following we will view  $\mathcal{W}$  as a regular directed graph, rather than an AND-OR graph. The test is given by the procedure CHECK-SUPERCYCLE( $\mathcal{W}$ ) below, which works as follows. We first find the maximal strongly connected components (MSCC's) of  $\mathcal{W}$ . If no nontrivial MSCC's exist, then  $\mathcal{W}$  is acyclic and so the  $I$ -system is supercycle-free by Proposition 4.4. Otherwise, we execute the following check for each local-state node  $t_k$  in  $\mathcal{W}$ . If the

check marks  $t_k$  as “safe”, this means that no transition by  $P_k$  that ends in state  $t_k$  can create a supercycle where one did not exist previously. If all local-state nodes in  $\mathcal{W}$  are marked as “safe”, then we conclude that no transition by any process in the  $I$ -system can create a supercycle. Given that all initial  $I$ -states are supercycle-free, this then implies that every reachable  $I$ -state is supercycle free, and so the  $I$ -system is deadlock-free. The check for  $t_k$  is as follows. If  $t_k$  does not occur in a nontrivial MSCC of  $\mathcal{W}$ , then, by Proposition 4.3,  $t_k$  cannot occur in any supercycle, so mark  $t_k$  as safe and terminate. Otherwise, invoke CHECK-STATE( $t_k, C$ ), where  $C$  is the nontrivial MSCC of  $\mathcal{W}$  in which  $t_k$  occurs. Our test is sound but not complete. If some  $t_k$  is not marked “safe”, then we have no information about the possibility of the occurrence of supercycles.

CHECK-SUPERCYCLE( $\mathcal{W}$ )

1. Find the maximal strongly connected components of  $\mathcal{W}$
2. **for** each MSCC  $C$  of  $\mathcal{W}$  that consists of a single node  
    **if** the node is a local-state node **then** mark it “safe”
3. **for** each MSCC  $C$  of  $\mathcal{W}$  that contains more than one node  
    **for** each local-state node  $s_i$  of  $C$ , invoke CHECK-STATE( $s_i, C$ )
4. **if** all local-state nodes in  $\mathcal{W}$  are marked “safe”, **then return** (“No supercycle possible”)  
    **else return** (“Inconclusive”)

CHECK-STATE( $t_k, C$ )

1. Construct a subgraph  $SC$  of  $C$  as follows.  
    Let  $SC$  initially be  $C$   
    Remove from  $SC$  every  $s_k$  such that  $s_k \in \text{reachable}(P_k) - \{t_k\}$   
    **repeat** until no more nodes to remove from  $SC$   
        **if**  $a_j$  is a node in  $SC$  with no outgoing edges in  $SC$  **then**  
            let  $s_j$  be the unique node such that  $s_j \rightarrow a_j \in SC$   
            remove  $s_j$  and  $a_j$  and their incident edges from  $SC$
2. Compute the maximal strongly connected components of  $SC$
3. **if**  $t_k$  is not in some MSCC of  $SC$  **then** mark  $t_k$  as “safe” and terminate.  
    **else** Let  $MC$  be the MSCC of  $SC$  containing  $t_k$
4. **for** all  $(s_j, a_j, t_k, a_k, s_\ell)$  such that  $s_j \rightarrow a_j \rightarrow t_k \rightarrow a_k \rightarrow s_\ell \in MC$   
    Let  $J = \{\{j, k\}, \{k, \ell\}\}$   
    **if** there exists a state  $s_J$  of  $M_J$  such that:  
         $s_J$  is reachable along a path in  $M_J$  that ends in a transition by  $P_k$ , and  
         $s_j \rightarrow a_j \rightarrow t_k \rightarrow a_k \rightarrow s_\ell \in W_J(s_J)$   
    **then** mark all the nodes and edges in  $s_j \rightarrow a_j \rightarrow t_k \rightarrow a_k \rightarrow s_\ell$
5. Remove from  $MC$  all nodes and edges within two hops from  $t_k$  (in either direction) that are unmarked. Call the resulting graph  $MC'$
6. Calculate the maximal strongly connected components of  $MC'$
7. **if**  $t_k$  does not lie in an MSCC of  $MC'$  **then** mark  $t_k$  as safe

The procedure CHECK-STATE( $t_k, C$ ) tests whether the wait-for edges involving  $t_k$  and two other local states  $s_j, s_\ell$  can all be generated *simultaneously* in the triple system consisting of processes  $j, k, \ell$ . This generation must also be by a *reachable* transition of process  $k$ . If this generation cannot occur in

all such triple systems, then the wait-for-edges in the vicinity of  $s_k$  cannot be present simultaneously in some global state in a combination that results in a supercycle containing  $t_k$ .

**Theorem 4.5** *If all local-state nodes in  $\mathcal{W}$  are marked as “safe,” then the  $I$ -system  $P^I$  is supercycle-free.*

**Proof:** The proof idea is, roughly, that if all the wait-chains involving a local-state node  $t_k$  are unreachable in the relevant  $J$ -machines, then  $t_k$  cannot be part of a supercycle. The detailed proof is as follows.

We establish the contrapositive of the statement of the theorem. Let  $M_I$  be the global-state transition diagram of  $P^I$ . Assume that there exists a reachable global state  $v$  of  $M_I$  such that  $W_I(v)$  contains a supercycle  $SC$ . Let  $SC'$  be a supercycle as given by Proposition 4.3 (there may be more than one such supercycle—we choose one arbitrarily).

Since  $v$  is a reachable global state, there exists at least one finite computation path  $\pi = s_0 \dots v$  ending in  $v$ . By assumption, all initial global states are supercycle-free. Also, a supercycle is stable by definition: once created, it persists forever, since none of the involved processes can make a transition and thereby change their state. Hence, there exist states  $s, t$  along  $\pi$  such that (1)  $SC'$  is not contained in  $W_I(s')$  for any state  $s'$  along the prefix  $s_0 \dots s$  of  $\pi$ , and (2)  $SC'$  is contained in  $W_I(t')$  for every state  $t'$  along the suffix  $t \dots v$  of  $\pi$ .

Let  $P_k$  be the process that executes the transition from  $s$  to  $t$  along  $\pi$ . Since  $SC'$  is strongly connected, we have:

$$t_j \longrightarrow a_j \longrightarrow t_k \longrightarrow a_k \longrightarrow t_\ell \in SC' \quad (\text{a})$$

where  $t_j = t \downarrow j$ ,  $t_k = t \downarrow k$ ,  $t_\ell = t \downarrow \ell$ ,  $a_j$  is an arbitrarily chosen  $P_j$ -move in  $SC'$ ,  $a_k$  is an arbitrarily chosen  $P_k$ -move in  $SC'$ ,  $(j, k) \in I$ ,  $(k, \ell) \in I$ , and  $j$  may or may not be equal to  $\ell$ .

Now  $SC'$  is a subgraph of  $W_I(t)$ , by construction, hence by (a):

$$t_j \longrightarrow a_j \longrightarrow t_k \longrightarrow a_k \longrightarrow t_\ell \in W_I(t). \quad (\text{b})$$

By Proposition 4.2,  $W_I(t)$  is a subgraph of  $\mathcal{W}$ . So, from (b):

$$t_j \longrightarrow a_j \longrightarrow t_k \longrightarrow a_k \longrightarrow t_\ell \in \mathcal{W}. \quad (\text{c})$$

Let  $J = \{(j, k), (k, \ell)\}$ , and let  $t_J = t \downarrow J$ . By [2, Proposition 6.5.4.1]<sup>2</sup> and (b), we have

$$t_j \longrightarrow a_j \longrightarrow t_k \longrightarrow a_k \longrightarrow t_\ell \in W_J(t_J). \quad (\text{d})$$

Now  $t$  is reachable in  $M_I$ . Hence by [2, Corollary 6.4.5],  $t_J$  is reachable in  $M_J$ . From this and (d), it follows that all the nodes in the wait-chain  $t_j \longrightarrow a_j \longrightarrow t_k \longrightarrow a_k \longrightarrow t_\ell$  are marked in line 5 of CHECK-STATE( $t_k, C$ ). Thus, none of these nodes are removed in line 6. Since  $a_j$  and  $a_k$  are arbitrarily chosen, it follows that all nodes of  $SC'$  that are within two hops of  $t_k$  in either direction are marked in line 4 of CHECK-STATE( $t_k, C$ ). Thus, no nodes of  $SC'$  are deleted in line 5.

<sup>2</sup>Technically, the results of [2] were established only for concurrent programs with similar processes. However, the specific results that we use here were established without resorting to the similarity assumption, and so they also hold for concurrent programs with dissimilar processes.



Hence,  $SC'$  is a subgraph of the graph  $MC'$  that results after all deletions in line 5 are performed. Since  $SC'$  is strongly-connected,  $SC'$  lies in an MSCC of  $MC'$ . Since  $t_k$  is a node of  $SC'$ , it follows from line 7 that  $t_k$  is not marked as safe. Hence the contrapositive is established.  $\square$

**Proposition 4.6** *Let  $N$  be the size of the largest  $I$ -process (number of local states plus number of  $I$ -moves). Then the size of  $\mathcal{W}$  (number of nodes and edges) is  $O(K^2 N^2)$ .*

**Proof:** Each process contributes  $O(N)$  nodes to  $\mathcal{W}$ , and there are  $K$  processes. Thus,  $\mathcal{W}$  contains  $O(KN)$  nodes. Hence,  $\mathcal{W}$  contains  $O(K^2 N^2)$  edges.  $\square$

**Theorem 4.7** *The running time of  $CHECK\text{-}SUPERCYCLE(\mathcal{W})$  is  $O(K^4 N^4)$ .*

**Proof:** Line 1 of  $CHECK\text{-}SUPERCYCLE(\mathcal{W})$  takes time  $O(K^2 N^2)$ , using the algorithm in [6, Chapter 22].

Lines 2 and 4 of  $CHECK\text{-}SUPERCYCLE(\mathcal{W})$  clearly take time  $O(K^2 N^2)$ ,

Line 3 of  $CHECK\text{-}SUPERCYCLE(\mathcal{W})$  invokes  $CHECK\text{-}STATE(t_k, C)$ . Line 1 of  $CHECK\text{-}STATE(t_k, C)$  takes time  $O(|C|) +$  the time to compute  $reachable(P_k)$ .  $O(|C|)$  is  $O(K^2 N^2)$  since  $C$  is a subgraph of  $\mathcal{W}$ .  $reachable(P_k)$  can be computed in time  $O(K^2 N^2)$  since there are  $O(K^2)$  pair-systems, each of size  $O(N^2)$ . Line 2 of  $CHECK\text{-}STATE(t_k, C)$  takes time  $O(K^2 N^2)$  since  $C$  is a subgraph of  $\mathcal{W}$ . Line 3 of  $CHECK\text{-}STATE(t_k, C)$  takes time  $O(K^2 N^2)$  using a straightforward graph search. Line 4 takes time  $O(K^3 N^3)$  since there are  $O(K^3)$  three-process systems, each of size  $O(N^3)$ . Searching for wait-chains with  $MC$  is clearly in  $O(K^3 N^3)$ . Line 5 takes time  $O(N^2 K^2)$ , since  $|MC|$  is  $O(N^2 K^2)$ . Line 6 takes time  $O(N^2 K^2)$ , since  $|MC'|$  is  $O(N^2 K^2)$ . Line 7 takes time  $O(N^2 K^2)$ , using a straightforward graph search. Thus, a single invocation of  $CHECK\text{-}STATE(t_k, C)$  runs in time  $O(K^3 N^3)$ .  $CHECK\text{-}STATE(t_k, C)$  is invoked at most  $O(KN)$  times, since  $\mathcal{W}$  contains  $O(KN)$  nodes, and  $CHECK\text{-}STATE(t_k, C)$  is invoked at most once for each local-state node  $t_k$  in  $\mathcal{W}$ . Hence the total running time of line 3 of  $CHECK\text{-}SUPERCYCLE(\mathcal{W})$  is  $O(K^4 N^4)$ .

Thus,  $CHECK\text{-}SUPERCYCLE(\mathcal{W})$  runs in time  $O(K^4 N^4)$ .  $\square$

It may be possible to improve the runtime complexity of the algorithm using more sophisticated graph search strategies. For example, for each three-process system, we could collect all the wait-chains together and search for them all at once within the global state-transition graph (GSTD) of the three-process system. Wait-chains that are found could then be marked appropriately for subsequent processing.

It is not too hard to verify that the global wait-for graph for the hierarchical resource allocation strategy that we discussed in Section 3 is acyclic. Indeed, a supercycle in a wait-for graph represents a cyclical waiting pattern between processes. However, a hierarchy establishes a total order between processes, and the transitions in the graph represent blocking conditions, which can occur only when moves of a process with a lower priority are blocked by a process with higher priority. Thus, waiting conditions form chains, and not cycles in the wait-for graph. In a more general situation, the requirement of total hierarchical order can be relaxed for a subset of resources. Clearly, in this case deadlock can occur, depending on the sets of resources that each process attempts to acquire and the order of requests. Our algorithm can efficiently detect deadlocks in these cases.

Now we turn to some more examples, for which we explicitly construct the wait-for graph and show the presence or absence of supercycles.

**Example 4.8** [Verification of deadlock freedom] We study a special case of resource allocation problem [16] that we presented in Section 3. In this system, there are two resources (we refer to them as priority queues) and the additional parameter is the set of priorities of processes for the queues. Consider an  $I$ -system where the processes are partitioned into 3 classes, and are accessing two priority queues  $R$  and  $Q$ . The first class of processes has the highest priority for  $R$ , and the second class of processes has the highest priority for  $Q$ . The skeletons of processes  $P_1$ ,  $P_2$ , and  $P_3$  that represent the three classes of processes are described in Figure 6.

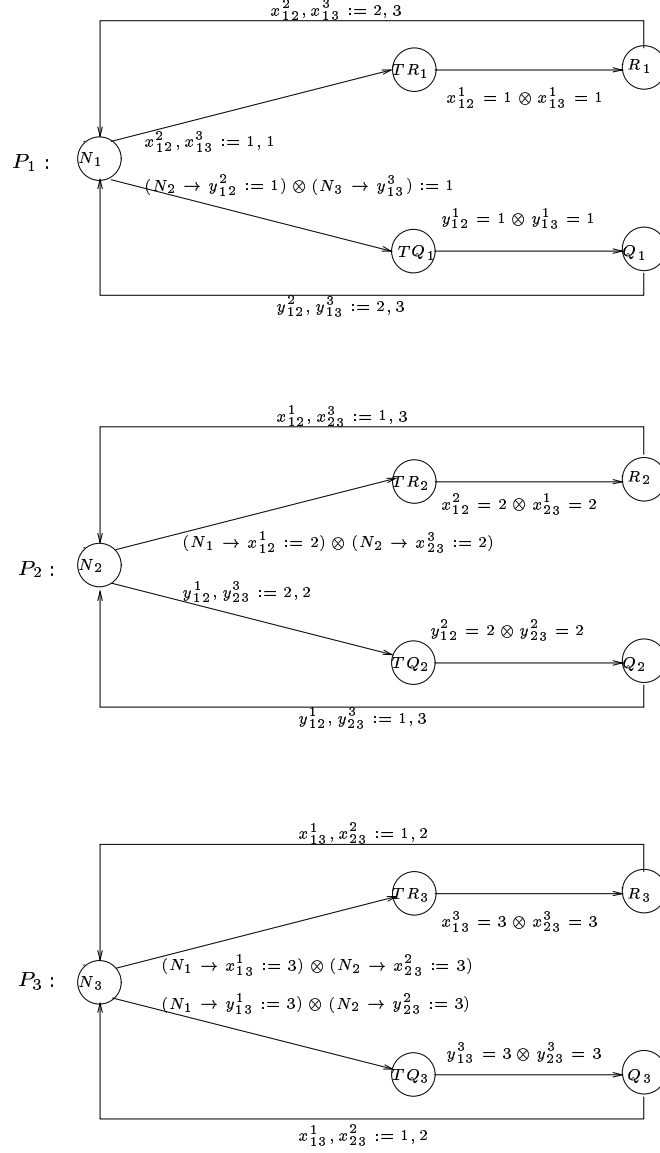


Figure 6: The example of three processes with two priority queues.

For processes in the same class and processes in different classes that have the same priority, the access to a queue is FIFO. There can be only one process at a time at the head of each queue (i.e., in

state  $R_i$  for queue  $R$  and  $Q_i$  for queue  $Q$ ). The system has a single initial state  $[N_1 N_2 N_3, \dots, N_K]$ . For clarity, we omit the shared variables that control the FIFO access to queues between the processes in the same class. Intuitively, a deadlock can occur if there are several processes with the same priority in a trying state ( $TR_i$  or  $TQ_i$ ). However, the guards on transitions to trying states guarantee that a process enters a trying state iff either there is no other process in the trying state, or the other process in the trying state has a lower priority.




Figure 7: A part of the global graph.

Figure 7 gives a fragment of  $\mathcal{W}$  for this system. The move nodes are labeled with the name of the source and target states and a directed edge between them (for example, the node that corresponds to the move from  $TR_1$  to  $R_1$  is labeled with  $TR_1 \rightarrow R_1$ ). For clarity, the graph shows only two processes  $P_i$  and  $P'_i$  from each class, for  $1 \leq i \leq 3$ , and only the trying states and their outgoing moves. It also omits the wait-for edges between the primed and unprimed processes from different classes. Still, it is easy to see that the graph contains many supercycles. We can prove that these supercycles are unreachable using the algorithm CHECK-STATE. Indeed, for all triples of processes that include an edge in a supercycle, all processes in the triple are in their trying states. These states are unreachable in the triple systems. Thus, the  $I$ -system passes our check, and hence is deadlock-free.

**Example 4.9** We now return to example 3.4. Recall, that the system in the example consists of 4 processes  $P_1, P_2, P_3$ , and  $P_4$  accessing two critical sections, where the processes  $P_3$  and  $P_4$  can block all other processes, and the processes  $P_1$  and  $P_2$  can only block each other. The test in Section 3 failed for this system in example despite it being deadlock-free. We sketch the application of the test in this section as follows. We apply CHECK-STATE to each state that participates in a supercycle. The only supercycles in the global wait-for graph are created by two processes - either  $P_1$  and  $P_2$ , or  $P_3$  and  $P_4$  in their trying states. By checking triples of processes, we can show that a trying state  $T_i$  and the assignment to shared variables that blocks two moves of another process are unreachable in all triples, for  $1 \leq i \leq 4$ . Thus, the system passes the test and is deadlock-free.

**Example 4.10** [A system with a deadlocked state] In this example we describe a system with a reachable deadlocked state and demonstrate the evidence for the deadlock in the the global wait-for graph. The system consists of two dissimilar processes  $P_1$  and  $P_2$  accessing two priority queues  $R$  and  $Q$ . The skeletons of the processes are shown in Figure 8. The system has a single initial state  $[N_1 N_2]$ . The process  $P_1$  has priority over  $P_2$  in the access to  $R$ , and the process  $P_2$  has priority over  $P_1$  in the access to  $Q$ . Both processes access both queues during one execution round. The skeletons of the processes are presented in Figure 8.





Figure 8: The example of two processes with two priority queues.

The global state  $[B_1 A_2]$  is reachable from the initial state  $[N_1 N_2]$  by the following path.

$$[N_1 N_2] \xrightarrow{1} [y_{12}^2 := 1, B_1 N_2] \xrightarrow{2} [x_{12}^1 := 2, B_1 A_2].$$

This state is deadlocked, since the process  $P_1$  is waiting for the process  $P_2$  to release  $Q$ , and the process  $P_2$  is waiting for the process  $P_1$  to release  $R$ . This cyclic waiting can be discovered by examining the global wait-for graph for supercycles. The drawing presents a fragment of the graph that contains the supercycle for the deadlocked state  $[B_1 A_2]$ .



## 5 Conclusions

We studied the problem of deadlock detection and avoidance in concurrent systems. We presented two sufficient conditions for assuring deadlock-freedom in a concurrent system. The first condition is suitable for systems of similar processes, while the second condition is suitable for systems of dissimilar processes. We presented efficient algorithms for computing both conditions. Our algorithms are polynomial in the number of processes and in the size of a single process.

## References

- [1] P. C. Attie. Synthesis of large concurrent programs via pairwise composition. In *CONCUR'99: 10th International Conference on Concurrency Theory*, number 1664 in LNCS, Aalborg, Denmark, Aug. 1999. Springer-Verlag.
- [2] P. C. Attie and E. A. Emerson. Synthesis of concurrent systems with many similar processes. *ACM Trans. Program. Lang. Syst.*, 20(1):51–115, Jan. 1998.
- [3] P.C. Attie. Synthesis of large dynamic concurrent programs from dynamic specifications. Technical report, Northeastern University, Boston, MA, 2003. Available at <http://www.ccs.neu.edu/home/attie/pubs.html>.
- [4] P.C. Attie. Finite-state concurrent programs can be expressed pairwise. Submitted, 2004.
- [5] E.G. Coffman, M.J. Elphick, and A. Shoshani. System deadlocks. *ACM Comput. Surv.*, 3:67–78, June 1971.
- [6] T.H. Cormen, C.E. Leiserson, R.L. Rivest, and C. Stein. *Introduction to Algorithms; Second Edition*. MIT Press and McGraw-Hill, 2001.
- [7] E. W. Dijkstra. *A Discipline of Programming*. Prentice-Hall Inc., Englewood Cliffs, N.J., 1976.
- [8] E. A. Emerson and E. M. Clarke. Using branching time temporal logic to synthesize synchronization skeletons. *Sci. Comput. Program.*, 2:241 – 266, 1982.
- [9] E. A. Emerson and A. P. Sistla. Symmetry and model checking. In *Proceedings of the 5th International Conference on Computer Aided Verification*, volume 697 of *Lecture Notes in Computer Science*, pages 463–477, Berlin, 1993. Springer-Verlag.
- [10] P. Godefroid. *Partial Order Methods for the Verification of Concurrent Systems*. PhD thesis, University of Liege, 1994.
- [11] P. Godefroid, D. Peled, and M. Staskauskas. Using partial-order methods in the formal validation of industrial concurrent programs. *Transactions on Software Engineering*, 22(7):496–507, Jan. 1996.
- [12] P. Godefroid and P. Wolper. A partial approach to model checking. *Information and Computation*, 110(2):305–326, 1991.
- [13] R. C. Holt. Some deadlock properties of computer systems. *ACM Computing Surveys*, 4(3):179–196, 1972.
- [14] E. Knapp. Deadlock detection in distributed databases. *ACM Comput. Surv.*, 19(4):303–328, Dec. 1987.
- [15] P. Ladkin and B. Simons. Compile-time analysis of communicating processes. In *Proc. International Conference on Supercomputing*, pages 248–259. ACM Press, 1992.
- [16] N. A. Lynch. *Distributed Algorithms*. Morgan-Kaufmann, San Francisco, California, USA, 1996.
- [17] D. Peled. Partial order reduction: Model-checking using representatives. In *MFCS*, 1996.

- [18] B. Rex. Inference of  $k$ -process behavior from two-process programs. Master's thesis, School of Computer Science, Florida International University, Miami, FL, April 1999.
- [19] A. S. Tanenbaum. *Modern Operating Systems, second edition*. Prentice-Hall, Englewood Cliffs, New Jersey, USA, 2001.